

Machine Learning (CE 40717)

Fall 2025

Ali Sharifi-Zarchi

CE Department
Sharif University of Technology

October 24, 2025



1 Introduction

2 Discriminant Functions

3 Linear Classifiers

4 Perceptron

5 Cost Functions

6 Imbalanced Data

7 Cross Validation


8 Multi-Category Classification

Definition

- **Given:** A training dataset

$$D = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^N, \quad y^{(i)} \in \{1, 2, \dots, K\}$$

- **Goal:** Learn a function that maps any new input \mathbf{x} to one of K classes.
- **Examples:**
 - Email \rightarrow Spam or Not Spam
 - Image \rightarrow Cat, Dog, or Bird
 - Transaction \rightarrow Fraudulent or Legitimate



Real-World Example of Classification

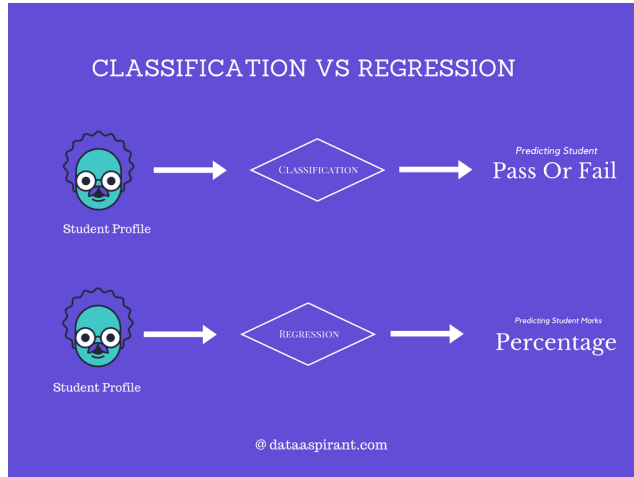
- **Pima Indians Diabetes Dataset:**
 - **Problem:** Predict whether a patient has diabetes based on medical diagnostics.
 - **Context:** Early detection of diabetes is critical for treatment and management.

	Number of times pregnant	Glucose	Blood Pressure	Skin Thickness	Insulin	Diabetes pedigree function	Age	BMI	Label
Patient 1	6	148	72	35	0	0.627	50	33.6	Positive
Patient 2	1	85	66	29	0	0.351	31	26.6	Negative
Patient 3	0	137	40	35	168	2.288	33	43.1	Positive
Patient 4	1	89	66	23	94	0.167	21	28.1	Negative
.
.
.

Classification vs. Regression: Comparison Table

Aspect	Regression	Classification
Output Type	Continuous value (\mathbb{R})	Discrete class label
Examples	House price, temperature	Spam detection, sentiment analysis
Evaluation Metrics	MSE, MAE	Accuracy, Precision, Recall

Classification vs. Regression: Figure



1 Introduction

② Discriminant Functions

③ Linear Classifiers

④ Perceptron

5 Cost Functions

6 Imbalanced Data

7 Cross Validation

9 Multi-Class Classification

Discriminant Functions in Machine Learning

- **Conceptual Overview:**

A discriminant function constitutes a mapping from the feature space to a real-valued score that quantifies the likelihood or confidence of a sample belonging to a specific class.

- **Formal Definition:**

Let $\mathbf{x} \in \mathbb{R}^d$ denote a feature vector. A discriminant function is a function $g(\mathbf{x}) : \mathbb{R}^d \rightarrow \mathbb{R}$ such that larger values of $g(\mathbf{x})$ correspond to stronger evidence for a particular class.

- **Objective:**

Design $g(\mathbf{x})$ to maximize correct classification over a given dataset.

Classification Using Discriminant Functions

- **Binary Classification:**

- Consider two discriminant functions $g_1(\mathbf{x})$ and $g_2(\mathbf{x})$ corresponding to classes C_1 and C_2 , respectively.
- The predicted class \hat{y} is determined by the criterion:

$$\hat{y} = \begin{cases} C_1 & \text{if } g_1(\mathbf{x}) > g_2(\mathbf{x}) \\ C_2 & \text{otherwise.} \end{cases}$$

- **Multi-Class Classification:**





- For k classes, compute $g_i(\mathbf{x})$ for each class C_i , $i = 1, \dots, k$.
- Assign \mathbf{x} to the class corresponding to the maximal discriminant value:

$$\hat{y} = \arg \max_i g_i(\mathbf{x})$$

- **Interpretation:** The discriminant function serves as a quantitative measure of class membership confidence.

Decision Boundary

- **Definition:** A dividing hyperplane that separates different classes in a feature space, also known as "Decision Surface".



Two-Class Discriminant Function

- In the case of a binary classification problem, a single discriminant function suffices:

$$g: \mathbb{R}^d \rightarrow \mathbb{R}.$$

- The two class-specific discriminants can be represented as:

$$g_1(\mathbf{x}) = g(\mathbf{x}), \quad g_2(\mathbf{x}) = -g(\mathbf{x}).$$

- The associated decision rule is:

$$\hat{y} = \begin{cases} C_1 & \text{if } g(\mathbf{x}) > 0 \\ C_2 & \text{if } g(\mathbf{x}) < 0 \end{cases}$$

- The decision boundary corresponds to the set $\{\mathbf{x} \mid g(\mathbf{x}) = 0\}$.
- This formulation provides a foundation for subsequent extension to multi-class discriminant analysis.

1 Introduction

2 Discriminant Functions

3 Linear Classifiers

4 Perceptron

5 Cost Functions

6 Imbalanced Data

7 Cross Validation

8 Multi-Category Classification

Linear Classifiers

- **Definition:**

Linear classifiers assign class labels using a decision function that is linear in the feature vector $\mathbf{x} \in \mathbb{R}^d$, or linear in a set of transformed features of \mathbf{x} .

- **Linearly separable data:**

Data points that can be perfectly separated by a linear decision boundary.

- **General form:**

$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0,$$

where \mathbf{w} defines the orientation of the decision surface and w_0 determines its position.

Two-Category Classification

- **Linear discriminant:**


$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

- $\mathbf{x} = [x_1, \dots, x_d]^T$, $\mathbf{w} = [w_1, \dots, w_d]^T$, w_0 : bias

- **Decision rule:**


$$\hat{y} = \begin{cases} C_1, & \text{if } g(\mathbf{x}) \geq 0 \\ C_2, & \text{otherwise} \end{cases}$$

- **Decision surface:** $\mathbf{w}^T \mathbf{x} + w_0 = 0$




Geometric Properties of Linear Decision Boundaries

- The decision boundary is a $(d - 1)$ -dimensional hyperplane in \mathbb{R}^d .
- **Properties:**
 - Orientation is determined by the normal vector $\mathbf{w}/\|\mathbf{w}\|$.
 - Bias w_0 controls the displacement along the normal vector.
- Points on opposite sides of the hyperplane are assigned to different classes.



Nonlinear Decision Boundaries

- **Problem:**
Many datasets cannot be separated by a linear hyperplane.
- **Feature Transformation:**
Map input vector \mathbf{x} to a higher-dimensional space $\phi(\mathbf{x})$.
- **Resulting Decision Boundary:**
Linear in the transformed space, but nonlinear in the original feature space.



1 Introduction

2 Discriminant Functions

3 Linear Classifiers

4 Perceptron

5 Cost Functions

6 Imbalanced Data

7 Cross Validation

8 Multi-Category Classification

What Is Perceptron?

- **Perceptron Unit:**

- **Basic Building Block:** The perceptron is the simplest form of an artificial neuron used for classification tasks.
- **Linear Classifier:** It computes a linear combination of input features followed by a threshold-based decision.
- **Binary Decision:** Outputs 1 if the weighted sum of inputs exceeds a threshold; otherwise outputs 0.
- **Components:** Consists of input features, associated weights, a bias term, and an activation function (commonly a step or sigmoid function).




Figure adapted from *Grokking Machine Learning*, L. G. Serrano.

Inspired by Biology

- **Biological Motivation Behind Perceptron:**
 - **Inspired by Neurons:** Perceptron mimics the basic function of biological neurons in the brain.
 - Input and Output, Activation Function.




Figure adapted from Grokking Machine Learning, L. G. Serrano.

Single Neuron

- **Single Neuron as a Linear Decision Boundary**

- **Mathematical Form:** The output of a single neuron is computed as:

$$y = f(\mathbf{w}^T \mathbf{x} + w_0)$$

where:

- \mathbf{x} is the input vector.
 - \mathbf{w} is the weight vector.
 - w_0 is the bias term.
 - f is an activation function (e.g., step function).
- **Linear Separation:** A neuron defines a linear decision boundary:
 $\mathbf{w}^T \mathbf{x} + w_0 = \text{threshold}$ (0 for step, 0.5 for sigmoid)
 - **Decision Rule:** C_1 if $\mathbf{w}^T \mathbf{x} + w_0 \geq \text{threshold}$, otherwise C_2 .

$$\text{Class} = f(\mathbf{w}^T \mathbf{x} + w_0)$$





Figure adapted from Grokking Machine Learning, L. G. Serrano.

Limitations of a Single Perceptron


- **What a Single Perceptron Can and Can't Do:**
 - **Performs Linear Separations:** A perceptron can handle linearly separable problems such as:
 - AND operation
 - OR operation
 - **Fails on Non-Linear Problems:** A single perceptron fails to solve non-linear problems like XOR, as the data points cannot be separated by a straight line.



Towards Complex Decision Boundaries

- **Multi-Layer Perceptron (MLP):**

- **Adding Layers for More Complexity:** An MLP consists of multiple layers of neurons that allow us to model more complex functions than a single neuron.
 - Each layer introduces new decision boundaries, making it possible to separate non-linear data.
- **Two-Layer Example:**
 - Input Layer → Hidden Layer → Output Layer
 - Hidden layer introduces non-linear transformations that enable complex decision regions.



Figures adapted from Grokking Machine Learning, L. G. Serrano.

Refining the Decision Boundary

- **New Neurons for Better Separation:** By adding more neurons to a layer, we can further refine the decision boundary to better separate complex data.
- Each additional neuron introduces new features that help the model make more accurate decisions.

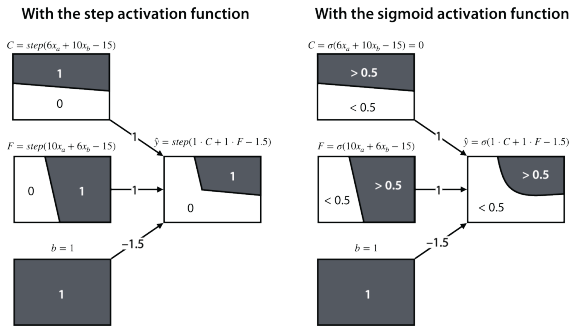


Figure adapted from Grokking Machine Learning, L. G. Serrano.

- 1 Introduction
- 2 Discriminant Functions
- 3 Linear Classifiers
- 4 Perceptron
- 5 Cost Functions**
- 6 Imbalanced Data
- 7 Cross Validation

Cost Functions

- **Understanding the Goal**

- In the perceptron, we use $\mathbf{w}^T \mathbf{x}$ to make predictions.
- Goal is to find the optimal \mathbf{w} so that the predicted labels match the true labels as much as possible.
- To achieve this, we define a cost function, which measures the **difference** between **predicted** and **actual** labels.
- Finding discriminant functions (\mathbf{w}^T, w_0) is framed as minimizing a cost function.
 - Based on training set $D = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^N$, a cost function $J(\mathbf{w})$ is defined.
 - Problem converts to finding optimal $\hat{g}(\mathbf{x}) = g(\mathbf{x}; \hat{\mathbf{w}})$ where


$$\hat{\mathbf{w}} = \arg \min_{\mathbf{w}} J(\mathbf{w})$$

Sum of Squared Error Cost Function

- Sum of Squared Error (SSE) Cost Function**

- Formula:** $J(\mathbf{w}) = \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2$, $\hat{y}^{(i)} = \mathbf{w}^T \mathbf{x}^{(i)} + w_0$
- SSE minimizes the magnitude of the error, which is ideal for regression but **irrelevant** for classification.
- If the model predicts close to the true class but not exactly 0 or 1, SSE still shows positive error, even for correct predictions.

- SSE is also prone to overfitting noisy data, as small variations can cause significant changes in the cost.



An Alternative for SSE Cost Function

- **Number of Misclassifications**

- **Definition:** Measures how many samples are misclassified by the model.
- **Formula:**

$$J(\mathbf{w}) = \sum_{i=1}^n \left(\frac{y^{(i)} - \text{sign}(\hat{y}^{(i)})}{2} \right)^2, \quad \hat{y}^{(i)} = \mathbf{w}^T \mathbf{x}^{(i)} + w_0, \quad y^{(i)} \in \{-1, +1\}$$

- **Limitations:**

- **Piecewise Constant:** The cost function is non-differentiable, so optimization techniques (like gradient descent) cannot be directly applied.

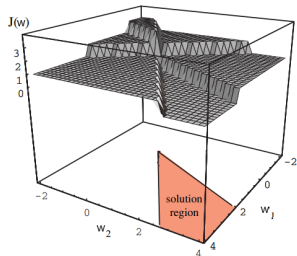



Figure adapted from Machine Learning and Pattern Recognition, Bishop

Perceptron Algorithm


- **The Perceptron Algorithm**

- **Purpose:** A simple algorithm for binary classification, separating two classes with a linear boundary.

$J(\mathbf{w})$



$J_P(\mathbf{w})$



Perceptron Criterion

- **Cost Function:** The perceptron criterion focuses on misclassified points:

$$J_p(\mathbf{w}) = - \sum_{i \in M} y^{(i)} \mathbf{w}^T \mathbf{x}^{(i)}, \quad y^{(i)} \in \{-1, +1\}$$

where M is the set of misclassified points.

- **Goal:** Minimize the loss by correctly classifying all points.

Batch Perceptron

- **Batch Perceptron:** Updates the weight vector using all misclassified points in each iteration.
- **Gradient Descent:** Adjusting weights in the direction that reduces the loss:

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} J_p(\mathbf{w})$$

$$\nabla_{\mathbf{w}} J_p(\mathbf{w}) = - \sum_{i \in M} y_i \mathbf{x}_i$$

- Batch Perceptron converges in finite number of steps for linearly separable data.

Single-sample Perceptron

- **Single Sample Perceptron:** Updates the weight vector after each individual point.
- **Stochastic Gradient Descent (SGD) Update Rule:**
 - Using only one misclassified sample at a time:

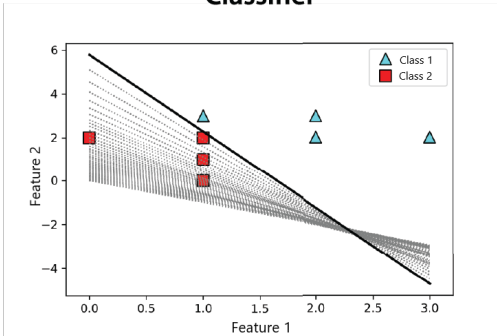
$$\mathbf{w} \leftarrow \mathbf{w} + \eta y_i \mathbf{x}_i$$

- Lower computational cost per iteration, faster convergence.
- If training data are linearly separable, the single-sample perceptron is also guaranteed to find a solution in a finite number of steps.


Example

- Perceptron changes \mathbf{w} in a direction that corrects error.

Classifier



Error



Figures adapted from Grokking Machine Learning, L. G. Serrano.

Convergence of the Perceptron — Theorem

Theorem: For linearly separable data with margin $\gamma > 0$ and $\|x_i\| \leq R$, the Perceptron algorithm makes at most $M \leq \frac{R^2}{\gamma^2}$ updates.

Notation:

- Dataset: $D = \{(x_i, y_i)\}_{i=1}^n$, with $x_i \in \mathbb{R}^d$, $y_i \in \{+1, -1\}$.
- Weight vector at step t : w_t , starting from $w_0 = 0$.
- Update rule (on mistake): $w_{t+1} = w_t + y_i x_i$.
- Assume there exists w^* with $\|w^*\| = 1$ that correctly classifies all samples.
- Each input is bounded: $\|x_i\| \leq R$ (after scaling, $R = 1$).
- Margin:

$$\gamma = \min_{(x_i, y_i) \in D} y_i (x_i^\top w^*) > 0.$$

Convergence of the Perceptron — Proof (1)

Let the algorithm make M mistakes.

- 1 Each mistake on (x_i, y_i) updates

$$w_{t+1} = w_t + y_i x_i.$$

- 2 By induction, $w_M = \sum_{t=1}^M y_t x_t$.

- 3 Inner product with w^* :

$$w_M \cdot w^* = \sum_{t=1}^M y_t (x_t \cdot w^*) \geq M\gamma.$$

Convergence of the Perceptron — Proof (2)

④ Norm growth:

$$\|w_M\|^2 = \|w_{M-1}\|^2 + 2y_M(w_{M-1} \cdot x_M) + \|x_M\|^2 \leq \|w_{M-1}\|^2 + R^2.$$

Hence, $\|w_M\| \leq R\sqrt{M}$.

⑤ Combine inequalities:

$$M\gamma \leq w_M \cdot w^* \leq \|w_M\| \|w^*\| \leq R\sqrt{M}.$$


⑥ Therefore,

$$M \leq \frac{R^2}{\gamma^2}.$$

Source: Novikoff (1962), *On Convergence Proofs for Perceptrons*; M. Collins, *Convergence Proof for the Perceptron Algorithm*, Columbia University.

Convergence of Perceptron Cont.

- **Non-Linearly Separable Data:** When no linear decision boundary can perfectly separate the classes, the Perceptron fails to converge.
 - If data is not linearly separable, there will always be some points that the model fails to classify.
 - As a result, the algorithm keeps adjusting the weights to fix the misclassified points, causing it to never converge.
 - For the data that are not linearly separable due to noise, **Pocket Algorithm** keeps in its pocket the best \mathbf{w} encountered up to now.



Pocket Algorithm

Algorithm 1 Pocket Algorithm

```

1: Initialize  $\mathbf{w}$ 
2: for  $t = 1$  to  $T$  do
3:    $i \leftarrow t \bmod N$ 
4:   if  $\mathbf{x}^{(i)}$  is misclassified then
5:      $\mathbf{w}^{new} = \mathbf{w} + \eta \mathbf{x}^{(i)} y^{(i)}$ 
6:     if  $E_{train}(\mathbf{w}^{new}) < E_{train}(\mathbf{w})$  then
7:        $\mathbf{w} = \mathbf{w}^{new}$ 
8:     end if
9:   end if
10: end for

```


$$\triangleright E_{train}(\mathbf{w}) = J_p(\mathbf{w})$$

7 Cross Validation

The Problem: Imbalanced Data

- Real-world datasets often contain classes with unequal representation.
- **Examples:**
 - Fraud detection: 0.1% Fraud, 99.9% Non-Fraud
 - Medical diagnosis: 2% Disease, 98% Healthy
- **Issue:** High accuracy may hide poor performance on the minority class.


Balanced vs. Imbalanced Datasets



Solution 1: Resampling Techniques

Undersampling

- Remove samples from majority class.
- + Reduces training time.
- – Risk of losing information.



Oversampling (e.g., SMOTE)

- Generate synthetic minority samples.
- + Preserves information.
- – May cause overfitting.



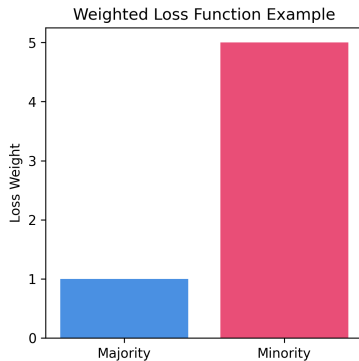
Solution 2: Algorithmic Approaches

Weighted Loss Function

- Assign a higher penalty to errors on the minority class.
- Encourages the model to focus on rare but important cases.

$$J(w) = - \sum_i w_{y^{(i)}} y^{(i)} \log(\hat{y}^{(i)})$$

where $w_{y^{(i)}}$ is larger for the minority class.



Hybrid Approaches

Combining Sampling Strategies

- Integrates both oversampling and undersampling to achieve better class balance.

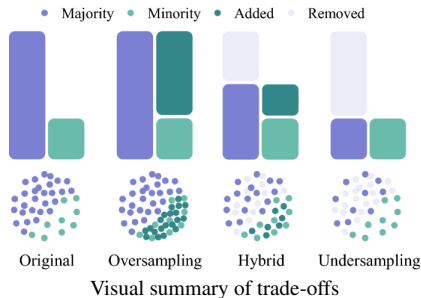
Typical Pipeline:

- ① Apply **SMOTE** to synthetically augment minority samples.
 - ② Randomly undersample the majority class to reduce imbalance.
 - ③ Train the model on the newly balanced dataset.
- Provides a practical trade-off between **bias** and **variance**.

Comparison of Methods

Summary

- **Sampling** — modifies data distribution
 (+ Simple, improves balance)
 (– Risk of over/underfitting)
- **Algorithmic** — adjusts model focus
 (+ No data change, interpretable)
 (– Needs careful weight tuning)
- **Hybrid** — combines both strategies
 (+ Balanced, strong results)
 (– More complex, slower training)



1 Introduction

② Discriminant Functions

③ Linear Classifiers

4 Perceptron

5 Cost Functions

6 Imbalanced Data

7 Cross Validation

9 Multi-Category Classification

Model Selection via Cross Validation

- **Cross-Validation**

- **Purpose:** Technique for evaluating how well a model generalizes to unseen data.
- **How It Works:** Split data into k folds; train on $k - 1$ folds and validate on the remaining fold.
- **Repeat Process:** Repeat k times, rotating the test fold each time. Average of all scores is the final score of the model.
- Cross-validation reduces overfitting and provides a more reliable estimation of model performance.

K-Fold Cross Validation

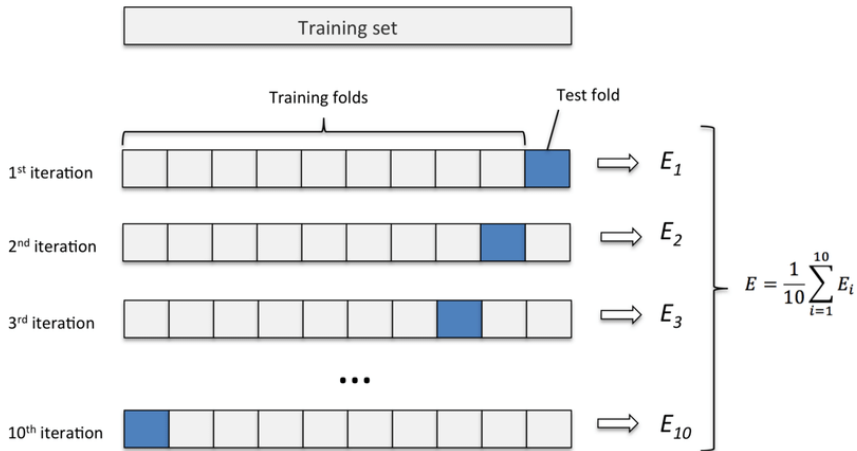



Figure adapted from Introduction to Support Vector Machines and Kernel Methods, J.M. Ashfaq.


Leave-One-Out Cross-Validation (LOOCV)

- **Leave-One-Out Cross-Validation (LOOCV)**
 - **How It Works:** Uses a single data point as the validation set ($k = 1$) and the rest as the training set. Repeat for all data points.
 - **Properties:**
 - **No Data Wastage:** Every data point is used for both training and validation.
 - **High Variance, Low Bias.**
 - **Computationally Expensive:** Requires training the model N times for N data points, making it slow for large datasets.
 - **Best for small datasets.**

Cross-Validation for Choosing Regularization Term



Cross-Validation for Choosing Model Complexity



9 Multi-Category Classification

Multi-Category Classification

- **Solutions to multi-category classification problem:**

- Extend the learning algorithm to support multi-class.
 - First, a function g_i for every class C_i is found.
 - Second, \mathbf{x} is assigned to C_i if $g_i(\mathbf{x}) > g_j(\mathbf{x}) \quad \forall i \neq j$

$$\hat{y} = \underset{i=1, \dots, c}{\operatorname{argmax}} g_i(\mathbf{x})$$

- Convert to a set of two-categorical problems.
 - Methods like **One-vs-Rest** or **One-vs-One**, where each classifier distinguishes between either **one class and the rest**, or **between pairs of classes**.

Multi-Category Classification: Linear Machines

- **Linear Machines:** Alternative to One-vs-Rest and One-vs-One methods; Each class is represented by its own discriminant function.

- **Decision Rule:**


$$\hat{y} = \underset{i=1,\dots,c}{\operatorname{argmax}} g_i(\mathbf{x})$$

The predicted class is the one with the highest discriminant function value.

- **Decision Boundary:** $g_i(\mathbf{x}) = g_j(\mathbf{x})$

$$(\mathbf{w}_i - \mathbf{w}_j)^T \mathbf{x} + (w_{0i} - w_{0j}) = 0$$

Linear Machines Cont.



- The decision regions of this discriminant are **convex** and **singly connected**. Any point on the line between two points within the same region can be expressed as $\mathbf{x} = \lambda \mathbf{x}_A + (1 - \lambda) \mathbf{x}_B$ where $\mathbf{x}_A, \mathbf{x}_B \in C_k$.

Multi-Class Perceptron Algorithm

- **Weight Vectors:**

- Maintain a weight matrix $W \in \mathbb{R}^{m \times K}$, where m is the number of features and K is the number of classes.
- Each column w_k of the matrix corresponds to the weight vector for class k .

$$\hat{y} = \operatorname{argmax}_{i=1, \dots, c} \mathbf{w}_i^T \mathbf{x}$$

$$J_p(\mathbf{W}) = - \sum_{i \in M} (\mathbf{w}_{y^{(i)}} - \mathbf{w}_{\hat{y}^{(i)}})^T \mathbf{x}^{(i)}$$

where M is the set of misclassified points.

Multi-Class Perceptron Algorithm

Algorithm 2 Multi-class perceptron

```
1: Initialize  $\mathbf{W} = [\mathbf{w}_1, \dots, \mathbf{w}_c], k \leftarrow 0$ 
2: while A pattern is misclassified do
3:    $k \leftarrow k + 1 \bmod N$ 
4:   if  $\mathbf{x}^{(i)}$  is misclassified then
5:      $\mathbf{w}_{\hat{y}^{(i)}} = \mathbf{w}_{\hat{y}^{(i)}} - \eta \mathbf{x}^{(i)}$ 
6:      $\mathbf{w}_{y^{(i)}} = \mathbf{w}_{y^{(i)}} + \eta \mathbf{x}^{(i)}$ 
7:   end if
8: end while
```

1 Introduction

② Discriminant Functions

③ Linear Classifiers

4 Perceptron

5 Cost Functions

6 Imbalanced Data

7 Cross Validation

9 Multi-Category Classification

Contributions

- **This slide has been prepared thanks to:**
 - Erfan Jafari
 - Aida Jalali

- [1] C. M. Bishop, *Pattern Recognition and Machine Learning*.
- [2] R. O. Duda, P. E. Hart, and D. G. Stork, *Pattern Classification*. 2001.
- [3] M. Soleymani, “Machine learning.” Sharif University of Technology.
- [4] S. F. S. Salehi, “Machine learning.” Sharif University of Technology.
- [5] Y. S. Abu-Mostafa, “Machine learning.” California Institute of Technology, 2012.
- [6] L. G. Serrano, *Grokking Machine Learning*. Manning Publications, 2020.
- [7] J. M. Ashfaq, “Introduction to support vector machines and kernel methods.” April 2019.