



- 1 Introduction
- 2 Logistic Regression
- 3 Extra reading
- 4 References

## 4 References

# Binary Classification Problem

- Consider a **binary classification** task:
  - Email classification: Spam / Not Spam
  - Online transactions: Fraudulent / Genuine
  - Tumor diagnosis: Malignant / Benign

Define the target variable formally:

$$y \in \{0, 1\}, \quad \begin{cases} 0 & \text{Negative class (e.g., benign tumor)} \\ 1 & \text{Positive class (e.g., malignant tumor)} \end{cases}$$

# Linear Regression for Classification


- A natural approach is to use linear regression:

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

and define a threshold at 0.5 for prediction:

$$\hat{y} = \begin{cases} 1, & h_{\theta}(x) \geq 0.5 \\ 0, & h_{\theta}(x) < 0.5 \end{cases}$$

- Here,  $\hat{y}$  is the **predicted class label** (i.e., the model's guess for  $y$ ). It converts the continuous output of  $h_{\theta}(x)$  into a discrete class (0 or 1).





## 1 Introduction

## ② Logistic Regression

## Fundamentals

## Decision surface

MLE and MAP

## Gradient descent

## Multi-class logistic regression

### 3 Extra reading

## 4 References

## 1 Introduction

## ② Logistic Regression

## Fundamentals

## Decision surface

MLE and MAP

## Gradient descent

## Multi-class logistic regression

### 3 Extra reading

## 4 References



# Introduction

- Suppose we have a binary classification task (so  $K = 2$ ).
- By observing **age**, **gender**, **height**, **weight** and **BMI** we try to distinguish if a person is **overweight** or **not overweight**.

| Age | Gender | Height (cm) | Weight (kg) | BMI  | Overweight |
|-----|--------|-------------|-------------|------|------------|
| 25  | Male   | 175         | 80          | 25.3 | 0          |
| 30  | Female | 160         | 60          | 22.5 | 0          |
| ... |        |             |             |      |            |
| 35  | Male   | 180         | 90          | 27.3 | 1          |

- We denote the **features** of a sample with vector  $x$  and the **label** with  $y$ .
- In logistic regression we try to find an  $\sigma(w^T x)$  which predicts **posterior** probabilities  $P(y = 1|x)$ .




## Introduction (cont.)

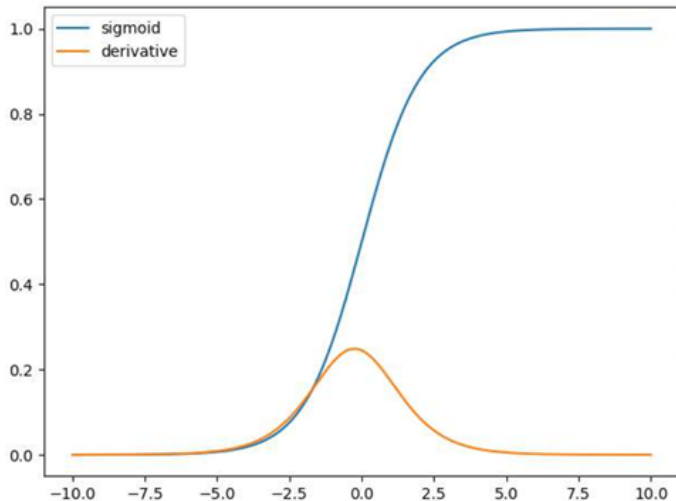
- Sigmoid (logistic) function.

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- A good candidate for activation function.
- It gives us a number between 0 and 1 **smoothly**.
- It is also **differentiable**



## Sigmoid function & its derivative





## 1 Introduction

## ② Logistic Regression

## Fundamentals

## Decision surface

MLE and MAP

## Gradient descent


## Multi-class logistic regression

### 3 Extra reading

## 4 References

# Decision surface

- Decision surface or decision boundary is the region of a problem space in which the output label of a classifier is ambiguous. (could be linear or non-linear)
- In binary classification it is where the probability of a sample belonging to each  $y = 0$  and  $y = 1$  is equal.



- Decision boundary hyperplane always has **one less dimension** than the feature space.

## Decision surface (cont.)

- An example of linear decision boundaries:


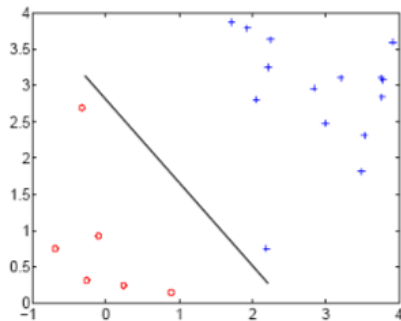


Figure adapted from Eric Xing, Machine Learning, CMU



## Decision surface (cont.)

- Back to our logistic regression problem.
- Decision surface  $\sigma(\mathbf{w}^T x) = \mathbf{constant}$ .



$$\sigma(\mathbf{w}^T x) = \frac{1}{1 + e^{-(\mathbf{w}^T x)}} = 0.5$$

- Decision surfaces are **linear functions** of  $x$ 
  - if  $\sigma(\mathbf{w}^T x) \geq 0.5$  then  $\hat{y} = 1$ , else  $\hat{y} = 0$
  - Equivalently, since  $\sigma(z) \geq 0.5$  only when  $z \geq 0$ , this means:
    - if  $\mathbf{w}^T x \geq 0$  then decide  $\hat{y} = 1$ , else  $\hat{y} = 0$

**$\hat{y}$  is the predicted label**

# Decision boundary example

$$\sigma(\mathbf{w}^T x) = \sigma(w_0 + w_1 x_1 + w_2 x_2)$$


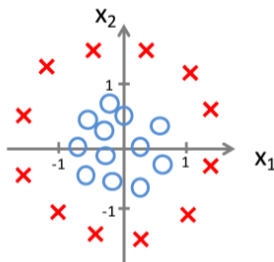


Predict  $y = 1$  if  $-3 + x_1 + x_2 \geq 0$

# Non-linear decision boundary example

$$\sigma(\mathbf{w}^T x) = \sigma(w_0 + w_1 x_1 + w_2 x_2 + w_3 x_1^2 + w_4 x_2^2)$$

We can learn more complex decision boundaries when having higher order terms



$$\text{Predict } y = 1 \text{ if } -1 + x_1^2 + x_2^2 \geq 0$$



# Maximum Likelihood Estimation (MLE)

- For  $n$  independent samples, the likelihood is:

$$L(\mathbf{w}) = \prod_{i=1}^n P(y^{(i)} | x^{(i)}, \mathbf{w})$$

- For binary classification ( $y \in \{0, 1\}$ ):

$$P(y^{(i)} | x^{(i)}, \mathbf{w}) = \sigma(\mathbf{w}^T x^{(i)})^{y^{(i)}} (1 - \sigma(\mathbf{w}^T x^{(i)}))^{1-y^{(i)}}$$

- This compact form works because  $y^{(i)} \in \{0, 1\}$ :
  - If  $y^{(i)} = 1$ , the expression becomes  $P(y = 1 | \dots) = \sigma(\mathbf{w}^T x^{(i)})$ .
  - If  $y^{(i)} = 0$ , the expression becomes  $P(y = 0 | \dots) = 1 - \sigma(\mathbf{w}^T x^{(i)})$ .
- Log-likelihood:

$$\ell(\mathbf{w}) = \sum_{i=1}^n [y^{(i)} \log \sigma(\mathbf{w}^T x^{(i)}) + (1 - y^{(i)}) \log(1 - \sigma(\mathbf{w}^T x^{(i)}))]$$

## From Likelihood to Cost Function

- Maximizing the likelihood  $\Leftrightarrow$  minimizing the negative log-likelihood (NLL):

$$J_{\text{MLE}}(\mathbf{w}) = -\ell(\mathbf{w})$$

- Can be written as an integral over the data distribution:

$$J_{\text{MLE}}(\mathbf{w}) = -\int p(x, y) \log P(y|x, \mathbf{w}) dx dy$$

- Empirical estimate (training data):

$$J_{\text{MLE}}(\mathbf{w}) = -\frac{1}{n} \sum_{i=1}^n \log P(y^{(i)}|x^{(i)}, \mathbf{w})$$

- This is exactly the **cross-entropy loss** used in classification.

# Maximum A Posteriori Estimation (MAP)

- MAP incorporates prior knowledge about parameters:

$$\hat{\mathbf{w}}_{\text{MAP}} = \arg \max_{\mathbf{w}} P(\mathbf{w}|D)$$

- Using Bayes' rule:

$$P(\mathbf{w}|D) \propto P(D|\mathbf{w})P(\mathbf{w})$$

- Equivalently:

$$\hat{\mathbf{w}}_{\text{MAP}} = \arg \max_{\mathbf{w}} \left[ \log P(D|\mathbf{w}) + \log P(\mathbf{w}) \right]$$

- Cost function:

$$J_{\text{MAP}}(\mathbf{w}) = - \sum_{i=1}^n \log P(y^{(i)}|x^{(i)}, \mathbf{w}) - \log P(\mathbf{w})$$

## MAP with Gaussian Prior (L2 Regularization)

- Gaussian prior:  $\mathbf{w} \sim \mathcal{N}(0, \tau^2 I)$

$$P(\mathbf{w}) \propto \exp\left(-\frac{\|\mathbf{w}\|^2}{2\tau^2}\right)$$

- MAP cost function:

$$J_{\text{MAP}}(\mathbf{w}) = J_{\text{MLE}}(\mathbf{w}) + \frac{1}{2\tau^2} \|\mathbf{w}\|^2$$

- Let  $\lambda = 1/\tau^2$ :

$$J_{\text{MAP}}(\mathbf{w}) = J_{\text{MLE}}(\mathbf{w}) + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- Equivalent to **L2-regularized logistic regression**.



# MAP with Laplace Prior (L1 Regularization)

- Laplace prior:  $\mathbf{w} \sim \text{Laplace}(0, b)$

$$P(\mathbf{w}) \propto \exp\left(-\frac{\|\mathbf{w}\|_1}{b}\right)$$

- MAP cost:

$$J_{\text{MAP}}(\mathbf{w}) = J_{\text{MLE}}(\mathbf{w}) + \lambda \|\mathbf{w}\|_1$$

where  $\lambda = 1/b$

- L1 penalty encourages **sparsity** (feature selection).

# Regularization Effects (L1 vs L2)

- **L2 (Ridge):** smooths weights, keeps all features but shrinks them.
- **L1 (Lasso):** drives some weights to zero, performs feature selection.
- Both prevent overfitting by penalizing model complexity.

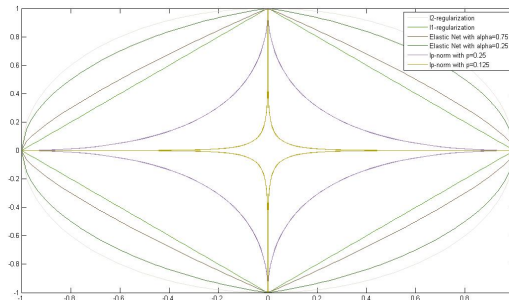


Image adapted from CS 4780, Cornell University

# Regularizers: Types and Properties

| Type | Form  | Properties / Advantages                         | Disadvantages / Effect              |
|------|---|---|-------------------------------------|
| L2   | $r(\mathbf{w}) = \ \mathbf{w}\ _2^2$                                      | Strictly convex, differentiable                 | Dense solutions (uses all features) |
| L1   | $r(\mathbf{w}) = \ \mathbf{w}\ _1$  | Convex, encourages sparsity                     | Not differentiable at 0             |
| Lp   | $r(\mathbf{w}) = \ \mathbf{w}\ _p = (\sum_i  v_i ^p)^{1/p}, 0 < p \leq 1$ | Very sparse solutions, initialization dependent | Non-convex, not differentiable      |

*Note:* Choice of regularizer affects sparsity and smoothness of learned weights.

# MLE vs MAP Comparison

|                     | <b>MLE</b>                            | <b>MAP</b>  |
|---------------------|---------------------------------------|---|
| Objective           | Maximize likelihood $P(D \mathbf{w})$ | Maximize posterior $P(D \mathbf{w})P(\mathbf{w})$ |
| Prior               | None (uniform)                        | Explicit prior $P(\mathbf{w})$                    |
| Cost function       | $-\log P(D \mathbf{w})$               | $-\log P(D \mathbf{w}) - \log P(\mathbf{w})$      |
| Regularization      | None                                  | Arises from prior                                 |
| Overfitting control | No                                    | Yes   |

## 1 Introduction

## 2 Logistic Regression

Fundamentals

Decision surface

MLE and MAP

**Gradient descent**

Multi-class logistic regression

## 3 Extra reading

## 4 References

# Gradient descent

- Remember from previous slides:

$$J(w) = \sum_{i=1}^n -y^{(i)} \log(\sigma(\mathbf{w}^T x^{(i)})) - (1 - y^{(i)}) \log(1 - \sigma(\mathbf{w}^T x^{(i)}))$$

- Update rule for **gradient descent**:

$$w^{t+1} = w^t - \eta \nabla_w J(w^t)$$

- With  $J(w)$  definition for logistic regression we get:

$$\nabla_w J(w) = \sum_{i=1}^n (\sigma(\mathbf{w}^T x^{(i)}) - y^{(i)}) x^{(i)}$$

# Gradient descent

- Compare the gradient of **logistic regression** with the gradient of **SSE** in **linear regression** :

$$\nabla_w J(w) = \sum_{i=1}^n (\sigma(\mathbf{w}^T x^{(i)}) - y^{(i)}) x^{(i)}$$

$$\nabla_w J(w) = \sum_{i=1}^n (\mathbf{w}^T x^{(i)} - y^{(i)}) x^{(i)}$$

## Loss function

- Loss function is a single overall measure of loss incurred for taking our decisions (over entire dataset).
- This is the **cross-entropy** (or log) loss for a single sample:

$$Loss(y, \sigma(\mathbf{w}^T x)) = -y \log(\sigma(\mathbf{w}^T x)) - (1 - y) \log(1 - \sigma(\mathbf{w}^T x))$$

- Since in binary classification  $y \in \{0, 1\}$ , the loss simplifies:

$$Loss(y, \sigma(\mathbf{w}^T x)) = \begin{cases} -\log(\sigma(\mathbf{w}^T x)) & \text{if } y = 1 \\ -\log(1 - \sigma(\mathbf{w}^T x)) & \text{if } y = 0 \end{cases}$$







## ② Logistic Regression

## Multi-class logistic regression

## 4 References

# Multi-class logistic regression

- Now consider a problem where we have  $K$  classes and every sample only belongs to one class (for simplicity).



## Multi-class logistic regression (cont.)

- For each class  $k$ ,  $\sigma_k(x; \mathbf{W})$  predicts the probability of  $y = k$ .
  - i.e.,  $P(y = k|x, \mathbf{W})$
- For each data point  $x_0$ ,  $\sum_{k=1}^K P(y = k|x_0, \mathbf{W})$  must be 1
  - $W$  denotes a matrix of  $w_i$ 's in which each  $w_i$  is a weight vector dedicated for class label  $i$ .
- On a new input  $x$ , to make a prediction, we pick the class that maximizes  $\sigma_k(x; \mathbf{W})$ :

$$\alpha(x) = \underset{k=1, \dots, K}{\operatorname{argmax}} \sigma_k(x; \mathbf{W})$$

**if  $\sigma_k(x; \mathbf{W}) > \sigma_j(x; \mathbf{W}) \forall j \neq k$  then decide  $C_k$**



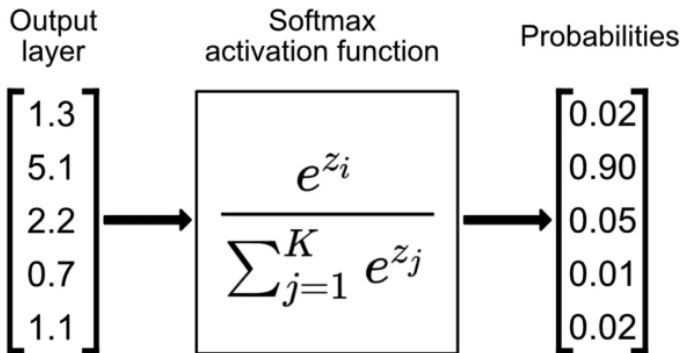
## Multi-class logistic regression (cont.)

- Softmax function **smoothly** highlights the maximum probability and is differentiable.
- Compare it with `max(.)` function which is strict and non-differentiable
- Softmax can also handle negative values because we are using exponential function
- And it gives us probability for each class since:

$$\sum_{k=1}^K \frac{\exp(w_k^T x)}{\sum_{j=1}^K \exp(w_j^T x)} = 1$$

## Multi-class logistic regression (cont.)

- An example of applying softmax (note that  $z_i = w^T x_i$ ):







## Multi-class logistic regression (cont.)

- From previous slides we have:

$$J(W) = - \sum_{i=1}^n \sum_{k=1}^K y_k^{(i)} \log(\sigma_k(x^{(i)}; \mathbf{W}))$$

- In which:

$$W = [w_1, w_2, \dots, w_K], \quad Y = \begin{pmatrix} y_1^{(1)} \\ y_1^{(2)} \\ \vdots \\ y_1^{(n)} \end{pmatrix} = \begin{pmatrix} y_1^{(1)} & \dots & y_K^{(1)} \\ y_1^{(2)} & \dots & y_K^{(2)} \\ \vdots & \ddots & \vdots \\ y_1^{(n)} & \dots & y_K^{(n)} \end{pmatrix}$$

- $y$  is a vector of length  $K$  (1-of- $K$  encoding)
  - For example  $y = [0, 0, 1, 0]^T$  when the target class is  $C_3$ .

## Multi-class logistic regression (cont.)

- Update rule for gradient descent:

$$w_j^{t+1} = w_j^t - \eta \nabla_w J(W^t)$$
$$\nabla_{w_j} J(W) = \sum_{i=1}^n (\sigma_j(x^{(i)}; \mathbf{W}) - y_j^{(i)}) x^{(i)}$$

- $w_j^t$  denotes the weight vector for class  $j$  (since in multi-class LR, each class has its own weight vector) in the  $t$ -th iteration

## 1 Introduction

## ② Logistic Regression

### 3 Extra reading

## Probabilistic view in classification

## Probabilistic classifiers

## 4 References

- 1 Introduction
- 2 Logistic Regression
- 3 **Extra reading**
  - Probabilistic view in classification
  - Probabilistic classifiers
- 4 References

## Probabilistic view in classification problem

- In a classification problem:
  - Each **feature** is a **random variable** (e.g. a person's height)
  - The **class label** is also considered a **random variable** (e.g. a person could be overweight or not)
- We observe the feature values for a random sample and intend to find its class label
  - Evidence: Feature vector  $x$
  - Objective: Class label

## Definitions

- Posterior probability : The probability of a class label  $C_k$  given a sample  $x$

$$P(C_k|x)$$

- Likelihood or class conditional probability : PDF of feature vector  $x$  for samples of class  $C_k$

$$P(x|C_k)$$

- Prior probability : Probability of the label be  $C_k$

$$P(C_k)$$

- $P(x)$ : PDF of feature vector  $x$ 
  - From total probability theorem:

$$P(x) = \sum_{k=1}^K P(x|C_k)P(C_k)$$

## 1 Introduction

## 2 Logistic Regression

### 3 Extra reading

## Probabilistic view in classification

## Probabilistic classifiers

## 4 References

## Probabilistic classifiers

- Probabilistic approaches can be divided in two main categories:
  - Generative
    - Estimate PDF  $P(x, C_k)$  for each class  $C_k$  and then use it to find  $P(C_k|x)$ . Alternatively estimate both PDF  $P(x|C_k)$  and  $P(C_k)$  to find  $P(C_k|x)$ .
  - Discriminative
    - Directly estimate  $P(C_k|x)$  for class  $C_k$



## Probabilistic classifiers (cont.)

- Let's assume we have input data  $x$  and want to classify the data into labels  $y$ .
- A generative model learns the **joint** probability distribution  $P(x, y)$ .
- A discriminative model learns the **conditional** probability distribution  $P(y|x)$

## Discriminative vs. Generative : example

- Suppose we have the following dataset in form of  $(x, y)$ :

 $(1, 0), (1, 0), (2, 0), (2, 1)$ 

- $P(x, y)$  is :

|         | $y = 0$       | $y = 1$       |
|---------|---------------|---------------|
| $x = 1$ | $\frac{1}{2}$ | 0             |
| $x = 2$ | $\frac{1}{4}$ | $\frac{1}{4}$ |

- $P(y|x)$  is :

|         | $y = 0$       | $y = 1$       |
|---------|---------------|---------------|
| $x = 1$ | 1             | 0             |
| $x = 2$ | $\frac{1}{2}$ | $\frac{1}{2}$ |

## Discriminative vs. Generative : example (cont.)

- The distribution  $P(y|x)$  is the natural distribution for classifying a given sample  $x$  into class  $y$ .
  - This is why that algorithms which model this directly are called **discriminative** algorithms.
- Generative algorithms model  $P(x, y)$ , which can be transformed into  $P(y|x)$  by Bayes rule and then used for classification.
  - However, the distribution  $P(x, y)$  can also be used for other purposes.
  - For example we can use  $P(x, y)$  to **generate** likely  $(x, y)$  pairs

## Generative approach

## 1 Inference

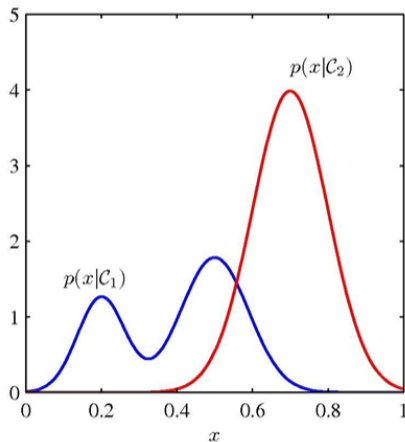
- Determine class conditional densities  $P(x|C_k)$  and priors  $P(C_k)$
- Use Bayes theorem to find  $P(C_k|x)$

## ② Decision

- Make optimal assignment for new input (after learning the model in the inference stage)
- if  $P(C_i|x) > P(C_j|x) \forall j \neq i$ , then decide  $C_i$ .

## Generative approach (cont.)

- Generative approach for a binary classification problem:



Figures adapted from Machine Learning and Pattern Recognition, Bishop

## Discriminative approach

## 1 Inference

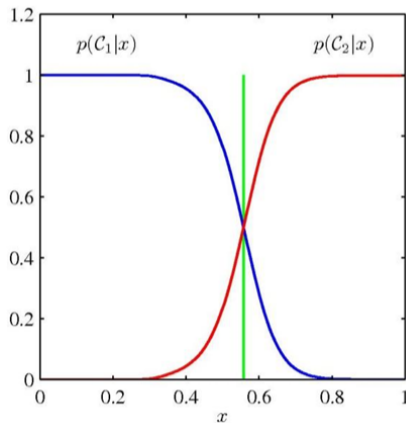
- Determine the posterior class probabilities  $P(C_k|x)$  directly.

## 2 Decision

- Make optimal assignment for new input (after learning the model in the inference stage)
- if  $P(C_i|x) > P(C_j|x) \forall j \neq i$ , then decide  $C_i$ .

## Discriminative approach (cont.)

- Discriminative approach for a binary classification problem:



Figures adapted from Machine Learning and Pattern Recognition, Bishop

- 1 Introduction
- 2 Logistic Regression
- 3 Extra reading
- 4 References



## Contributions

- **These slides are authored by:**
  - Danial Gharib

- [1] M. Soleymani Baghshah, “Machine learning.” Lecture slides.
- [2] A. Ng, “Ml-005, lecture 6.” Lecture slides.
- [3] C. M. Bishop, *Pattern Recognition and Machine Learning*. Information Science and Statistics, New York, NY: Springer, 1 ed., Aug. 2006.
- [4] S. Fidler, “Csc411.” Lecture slides.
- [5] A. Ng and T. Ma, *CS229 Lecture Notes*. Updated June 11, 2023.