

# Machine Learning (CE 40717)

Fall 2025

Ali Sharifi-Zarchi

CE Department  
Sharif University of Technology

November 1, 2025



- ① Introduction
- ② Decision surface
- ③ MLE and MAP
- ④ Gradient descent
- ⑤ Multi-class logistic regression
- ⑥ References

- 1 Introduction
- 2 Decision surface
- 3 MLE and MAP
- 4 Gradient descent
- 5 Multi-class logistic regression
- 6 References

# Binary Classification Problem

- Consider a **binary classification** task:
  - Email classification: Spam / Not Spam
  - Online transactions: Fraudulent / Genuine
  - Tumor diagnosis: Malignant / Benign

Define the target variable formally:

$$y \in \{0, 1\}, \quad \begin{cases} 0 & \text{Negative class (e.g., benign tumor)} \\ 1 & \text{Positive class (e.g., malignant tumor)} \end{cases}$$

# Linear Regression for Classification


- A natural approach is to use linear regression:

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

and define a threshold at 0.5 for prediction:


$$\hat{y} = \begin{cases} 1, & h_{\theta}(\mathbf{x}) \geq 0.5 \\ 0, & h_{\theta}(\mathbf{x}) < 0.5 \end{cases}$$

- Here,  $\hat{y}$  is the **predicted class label** (i.e., the model's guess for  $y$ ). It converts the continuous output of  $h_{\theta}(\mathbf{x})$  into a discrete class (0 or 1).



# Limitations of Linear Regression for Classification

- The model's output,  $h_{\theta}(\mathbf{x})$ , is unbounded and can produce predictions greater than 1 or less than 0.
- Linear regression does not provide probabilistic outputs.
- The decision boundary may be highly sensitive to outliers.



Requirement:  $0 \leq h_{\theta}(\mathbf{x}) \leq 1$

# Limitations of Perceptron with Step Activation for Classification

- As we saw earlier, perceptron Uses a step activation

$$f(z) = \begin{cases} 1, & z \geq 0 \\ 0, & z < 0 \end{cases}$$

- The step function has **slope = 0** almost everywhere (non-differentiable at 0).
- No gradient means it cannot use gradient-based optimization.
- Produces hard class labels (0/1), not probabilities, thus we cannot measure model confidence — probabilistic outputs are preferred.

# Properties of a Good Classifier

Therefore, we prefer a model that:

- Outputs **probabilities** for each class, so it has **bounded outputs** in  $[0, 1]$ , making it less sensitive to outliers.
- **Shows model confidence** through its probabilistic output.
- Has a **differentiable, convex objective with non-zero derivative**, allowing efficient optimization with gradient-based methods.



# Introduction

- Suppose we have a binary classification task (so  $K = 2$ ).
- By observing **age**, **gender**, **height**, **weight** and **BMI** we try to distinguish if a person is **overweight** or **not overweight**.

Age	Gender	Height (cm)	Weight (kg)	BMI	Overweight
25	Male	175	80	25.3	0
30	Female	160	60	22.5	0
...					
35	Male	180	90	27.3	1

- We denote the **features** of a sample with vector  $\mathbf{x}$  and the **label** with  $y$ .
- In logistic regression we try to find an  $\sigma(\mathbf{w}^\top \mathbf{x})$  which predicts **posterior** probabilities  $\mathbb{P}(y = 1 | \mathbf{x})$ .

# Introduction (cont.)

- $\sigma(\mathbf{w}^\top x)$ : probability that  $y = 1$  given  $x$  (parameterized by  $\mathbf{w}$ )

$$\mathbb{P}(y = 1 | \mathbf{x}, \mathbf{w}) = \sigma(\mathbf{w}^\top x)$$

$$\mathbb{P}(y = 0 | \mathbf{x}, \mathbf{w}) = 1 - \sigma(\mathbf{w}^\top x)$$

- We need to look for a function which gives us an output in the range  $[0, 1]$ . (like a probability).
- Let's denote this function with  $\sigma(\cdot)$  and call it the **activation function**.



## Introduction (cont.)

- Sigmoid (logistic) function:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- Smooth, bounded in  $(0, 1)$ , and **differentiable**.

$$\begin{aligned}\sigma'(z) &= \frac{e^{-z}}{(1+e^{-z})^2} = \underbrace{\frac{1}{1+e^{-z}}}_{\sigma(z)} \underbrace{\frac{e^{-z}}{1+e^{-z}}}_{1-\sigma(z)} \\ &= \sigma(z)(1-\sigma(z))\end{aligned}$$



## Introduction (cont.)

- The sigmoid function takes a number as input but we have:

$$\mathbf{x} = [x_0 = 1, x_1, \dots, x_d]$$


$$\mathbf{w} = [w_0, w_1, \dots, w_d]$$

- So we can use the **dot product** of  $\mathbf{x}$  and  $\mathbf{w}$ .
- We have  $0 \leq \sigma(\mathbf{w}^\top \mathbf{x}) \leq 1$ . which is the estimated probability of  $y = 1$  on input  $x$ .
- An Example : A basketball game (Win, Lose)
  - $\sigma(\mathbf{w}^\top \mathbf{x}) = 0.7$
  - In other terms 70 percent chance of winning the game.

- ① Introduction
- ② Decision surface
- ③ MLE and MAP
- ④ Gradient descent
- ⑤ Multi-class logistic regression
- ⑥ References

# Decision surface

- Decision surface or decision boundary is the region of a problem space in which the output label of a classifier is ambiguous. (could be linear or non-linear)
- In binary classification it is where the probability of a sample belonging to each  $y = 0$  and  $y = 1$  is equal.



- Decision boundary hyperplane always has **one less dimension** than the feature space.

# Decision surface (cont.)

- An example of linear decision boundaries:




Figure adapted from Eric Xing, Machine Learning, CMU

## Decision surface (cont.)

- Back to our logistic regression problem.
- Decision surface  $\sigma(\mathbf{w}^\top \mathbf{x}) = \mathbf{constant}$ .

$$\sigma(\mathbf{w}^\top \mathbf{x}) = \frac{1}{1 + e^{-(\mathbf{w}^\top \mathbf{x})}} = 0.5$$



- Decision surfaces are **linear functions** of  $x$ 
  - if  $\sigma(\mathbf{w}^\top \mathbf{x}) \geq 0.5$  then  $\hat{y} = 1$ , else  $\hat{y} = 0$
  - Equivalently, since  $\sigma(z) \geq 0.5$  only when  $z \geq 0$ , this means:
    - if  $\mathbf{w}^\top \mathbf{x} \geq 0$  then decide  $\hat{y} = 1$ , else  $\hat{y} = 0$

**$\hat{y}$  is the predicted label**



# Decision boundary example

$$\sigma(\mathbf{w}^\top \mathbf{x}) = \sigma(w_0 + w_1 x_1 + w_2 x_2)$$



Predict  $y = 1$  if  $-3 + x_1 + x_2 \geq 0$



- 1 Introduction
- 2 Decision surface
- 3 MLE and MAP**
- 4 Gradient descent
- 5 Multi-class logistic regression
- 6 References

# Maximum Likelihood Estimation (MLE)

- For  $n$  independent samples, the likelihood is:

$$\mathcal{L}(\mathbf{w}) = \prod_{i=1}^n \mathbb{P}(y^{(i)} | \mathbf{x}^{(i)}, \mathbf{w})$$

- For binary classification ( $y \in \{0, 1\}$ ):

$$\mathbb{P}(y^{(i)} | \mathbf{x}^{(i)}, \mathbf{w}) = \sigma(\mathbf{w}^\top \mathbf{x}^{(i)})^{y^{(i)}} (1 - \sigma(\mathbf{w}^\top \mathbf{x}^{(i)}))^{1-y^{(i)}}$$

- This compact form works because  $y^{(i)} \in \{0, 1\}$ :
  - If  $y^{(i)} = 1$ , the expression becomes  $P(y = 1 | \mathbf{x}^{(i)}, \mathbf{w}) = \sigma(\mathbf{w}^\top \mathbf{x}^{(i)})$ .
  - If  $y^{(i)} = 0$ , the expression becomes  $P(y = 0 | \mathbf{x}^{(i)}, \mathbf{w}) = 1 - \sigma(\mathbf{w}^\top \mathbf{x}^{(i)})$ .
- Log-likelihood:

$$\log \mathcal{L}(\mathbf{w}) = \sum_{i=1}^n [y^{(i)} \log \sigma(\mathbf{w}^\top \mathbf{x}^{(i)}) + (1 - y^{(i)}) \log(1 - \sigma(\mathbf{w}^\top \mathbf{x}^{(i)}))]$$

# From Likelihood to Cost Function

- Maximizing the likelihood  $\Leftrightarrow$  minimizing the negative log-likelihood (NLL):

$$J_{\text{MLE}}(\mathbf{w}) = -\log \mathcal{L}(\mathbf{w})$$

- Can be written as an integral over the data distribution:

$$J_{\text{MLE}}(\mathbf{w}) = -\int \mathbb{P}(\mathbf{x}, y) \log \mathbb{P}(y|\mathbf{x}, \mathbf{w}) \, d\mathbf{x} \, dy$$

- Empirical estimate (training data):

$$J_{\text{MLE}}(\mathbf{w}) = -\frac{1}{n} \sum_{i=1}^n \log \mathbb{P}(y^{(i)}|\mathbf{x}^{(i)}, \mathbf{w})$$

- This is exactly the **cross-entropy loss** used in classification.

# Maximum A Posteriori Estimation (MAP)

- MAP incorporates prior knowledge about parameters:

$$\mathbf{w}_{\text{MAP}} = \arg \max_{\mathbf{w}} \mathbb{P}(\mathbf{w}|D)$$

- Using Bayes' rule:

$$\mathbb{P}(\mathbf{w}|D) \propto P(D|\mathbf{w})P(\mathbf{w})$$

- Equivalently:

$$\mathbf{w}_{\text{MAP}} = \arg \max_{\mathbf{w}} \left[ \log \mathbb{P}(D|\mathbf{w}) + \log \mathbb{P}(\mathbf{w}) \right]$$

- Cost function:

$$J_{\text{MAP}}(\mathbf{w}) = - \sum_{i=1}^n \log \mathbb{P}(y^{(i)}|\mathbf{x}^{(i)}, \mathbf{w}) - \log \mathbb{P}(\mathbf{w})$$

# MAP with Gaussian Prior (L2 Regularization)

- Gaussian prior:  $\mathbf{w} \sim \mathcal{N}(0, \sigma^2 \mathbf{I})$

$$\mathbb{P}(\mathbf{w}) \propto \exp\left(-\frac{\|\mathbf{w}\|_2^2}{2\sigma^2}\right)$$

- MAP cost:

$$J_{\text{MAP}}(\mathbf{w}) = J_{\text{MLE}}(\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2$$

- L2 penalty encourages **weight shrinkage** (smooth regularization, no sparsity).

# MAP with Laplace Prior (L1 Regularization)

- Laplace prior:  $\mathbf{w} \sim \text{Laplace}(0, b)$

$$\mathbb{P}(\mathbf{w}) \propto \exp\left(-\frac{\|\mathbf{w}\|_1}{b}\right)$$

- MAP cost:

$$J_{\text{MAP}}(\mathbf{w}) = J_{\text{MLE}}(\mathbf{w}) + \lambda \|\mathbf{w}\|_1$$

- L1 penalty encourages **sparsity** (feature selection).



# Ridge vs. Lasso (Recap)

Ridge Regression (L2)	Lasso Regression (L1)
Constraint: $\ \mathbf{w}\ _2^2 \leq t$ (L2 ball).	Constraint: $\ \mathbf{w}\ _1 \leq t$ (L1 ball).
Contours of $J(\mathbf{w})$ typically touch the constraint boundary at smooth points.	Contours often hit the sharp <b>corners</b> of the diamond.
Produces small but nonzero coefficients (no exact sparsity).	Produces <b>sparse</b> solutions (some coefficients exactly zero).
Differentiable everywhere; the regularizer $\ \mathbf{w}\ _2^2$ is smooth.	Not differentiable at $w_i = 0$ ; $\ \mathbf{w}\ _1$ has sharp corners.
Strongly convex $\Rightarrow$ guarantees a <b>unique optimum</b> .	Not strongly convex $\Rightarrow$ may yield <b>multiple optima</b> .

- ① Introduction
- ② Decision surface
- ③ MLE and MAP
- ④ Gradient descent**
- ⑤ Multi-class logistic regression
- ⑥ References

# Gradient descent

- Remember from previous slides:

$$J(\mathbf{w}) = \sum_{i=1}^n -y^{(i)} \log(\sigma(\mathbf{w}^\top \mathbf{x}^{(i)})) - (1 - y^{(i)}) \log(1 - \sigma(\mathbf{w}^\top \mathbf{x}^{(i)}))$$

- Update rule for **gradient descent**:

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \nabla_{\mathbf{w}} J(\mathbf{w}^t)$$

- With  $J(\mathbf{w})$  definition for logistic regression we get:

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \sum_{i=1}^n (\sigma(\mathbf{w}^\top \mathbf{x}^{(i)}) - y^{(i)}) \mathbf{x}^{(i)}$$

## Gradient descent (cont.)

Let's find the derivative from the previous slide step by step:

- For one sample, we have:

$$J_i = -y^{(i)} \log \sigma(z^{(i)}) - (1 - y^{(i)}) \log(1 - \sigma(z^{(i)})), \quad z^{(i)} = \mathbf{w}^\top \mathbf{x}^{(i)}$$

- Differentiate w.r.t.  $z^{(i)}$ :

$$\frac{\partial J_i}{\partial z^{(i)}} = -\frac{y^{(i)}}{\sigma(z^{(i)})} \sigma'(z^{(i)}) + \frac{1 - y^{(i)}}{1 - \sigma(z^{(i)})} \sigma'(z^{(i)}) = \sigma(z^{(i)}) - y^{(i)}$$

- By chain rule we get:

$$\frac{\partial J_i}{\partial \mathbf{w}} = \frac{\partial J_i}{\partial z^{(i)}} \cdot \frac{\partial z^{(i)}}{\partial \mathbf{w}} = (\sigma(z^{(i)}) - y^{(i)}) \mathbf{x}^{(i)}$$

- Summing over all samples gives the result:

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \sum_{i=1}^n (\sigma(\mathbf{w}^\top \mathbf{x}^{(i)}) - y^{(i)}) \mathbf{x}^{(i)}$$

## Gradient descent (cont.)

- Compare the gradient of **logistic regression** with the gradient of **SSE** in **linear regression** :

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \sum_{i=1}^n (\sigma(\mathbf{w}^\top \mathbf{x}^{(i)}) - y^{(i)}) \mathbf{x}^{(i)}$$

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \sum_{i=1}^n (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$$

# Loss function

- Loss function is a single overall measure of loss incurred for taking our decisions (over entire dataset).
- This is the **cross-entropy** (or log) loss for a single sample:

$$J(y, \sigma(\mathbf{w}^\top \mathbf{x})) = -y \log(\sigma(\mathbf{w}^\top \mathbf{x})) - (1 - y) \log(1 - \sigma(\mathbf{w}^\top \mathbf{x}))$$

- Since in binary classification  $y \in \{0, 1\}$ , the loss simplifies:

$$J(y, \sigma(\mathbf{w}^\top \mathbf{x})) = \begin{cases} -\log(\sigma(\mathbf{w}^\top \mathbf{x})) & \text{if } y = 1 \\ -\log(1 - \sigma(\mathbf{w}^\top \mathbf{x})) & \text{if } y = 0 \end{cases}$$

# Convexity of Cross-Entropy Loss

This is the logistic regression loss over the dataset:

$$J(\mathbf{w}) = \sum_{i=1}^n \left[ -y^{(i)} \log \sigma(\mathbf{w}^\top \mathbf{x}^{(i)}) - (1 - y^{(i)}) \log(1 - \sigma(\mathbf{w}^\top \mathbf{x}^{(i)})) \right],$$

Define  $\hat{y}^{(i)} = \sigma(\mathbf{w}^\top \mathbf{x}^{(i)})$ . Then the gradient is:

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \sum_{i=1}^n (\hat{y}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}.$$

The Hessian (second derivative matrix) is:

$$\nabla_{\mathbf{w}}^2 J(\mathbf{w}) = \sum_{i=1}^n \hat{y}^{(i)} (1 - \hat{y}^{(i)}) \mathbf{x}^{(i)} (\mathbf{x}^{(i)})^\top = \mathbf{X}^\top \mathbf{D} \mathbf{X},$$

where  $\mathbf{D} = \text{diag}(\hat{y}^{(i)} (1 - \hat{y}^{(i)}))$  and each  $\hat{y}^{(i)} (1 - \hat{y}^{(i)}) \geq 0$ . For any vector  $\mathbf{v}$ ,

$$\mathbf{v}^\top \nabla_{\mathbf{w}}^2 J(\mathbf{w}) \mathbf{v} = \sum_{i=1}^n \hat{y}^{(i)} (1 - \hat{y}^{(i)}) (\mathbf{v}^\top \mathbf{x}^{(i)})^2 \geq 0.$$

Hence, the Hessian is **positive semidefinite (PSD)**, and  $J(\mathbf{w})$  is **convex**.

## Loss function (cont.)

- Note that this is different from the **zero-one loss**, which simply counts misclassifications:

$$J_{0-1}(y, \hat{y}) = \begin{cases} 1 & \text{if } y \neq \hat{y} \\ 0 & \text{if } y = \hat{y} \end{cases}$$

( $\hat{y}$  is the predicted label and  $y$  is the true label)



- We use cross-entropy (logistic loss) instead of zero-one loss because the zero-one loss function is **non-differentiable** and **non-convex**.
- The cross-entropy loss is a smooth, convex, and differentiable substitute, which allows us to use optimization methods like gradient descent.



- 1 Introduction
- 2 Decision surface
- 3 MLE and MAP
- 4 Gradient descent
- 5 Multi-class logistic regression**
- 6 References

# Multi-class logistic regression

- Now consider a problem where we have  $K$  classes and every sample only belongs to one class (for simplicity).



# Multi-class logistic regression (cont.)

- For each class  $k$ ,  $\sigma_k(\mathbf{x}; \mathbf{W})$  predicts the probability of  $y = k$ .
  - i.e.,  $\mathbb{P}(y = k | \mathbf{x}, \mathbf{W})$
- For each data point  $\mathbf{x}$ ,  $\sum_{k=1}^K \mathbb{P}(y = k | \mathbf{x}, \mathbf{W})$  must be 1
  - $\mathbf{W}$  denotes a matrix of  $\mathbf{w}_i$ 's in which each  $\mathbf{w}_i$  is a weight vector dedicated for class label  $i$ .
- On a new input  $x$ , to make a prediction, we pick the class that maximizes  $\sigma_k(\mathbf{x}; \mathbf{W})$ :

$$\alpha(\mathbf{x}) = \underset{k=1, \dots, K}{\operatorname{argmax}} \sigma_k(\mathbf{x}; \mathbf{W})$$

**if  $\sigma_k(\mathbf{x}; \mathbf{W}) > \sigma_j(\mathbf{x}; \mathbf{W}) \ \forall j \neq k$  then decide  $C_k$**

# Multi-class logistic regression (cont.)

- $K > 2$  and  $y \in \{1, 2, \dots, K\}$

$$\sigma_k(\mathbf{x}, \mathbf{W}) = \mathbb{P}(y = k | \mathbf{x}) = \frac{\exp(\mathbf{w}_k^\top \mathbf{x})}{\sum_{j=1}^K \exp(\mathbf{w}_j^\top \mathbf{x})}$$

- Normalized exponential (Aka **Softmax**)
- if  $\mathbf{w}_k^\top \mathbf{x} \gg \mathbf{w}_j^\top \mathbf{x}$  for all  $j \neq k$  then  $\mathbb{P}(C_k | \mathbf{x}) \approx 1$  and  $\mathbb{P}(C_j | \mathbf{x}) \approx 0$
- Note : remember from Bayes theorem:

$$\mathbb{P}(C_k | \mathbf{x}) = \frac{\mathbb{P}(\mathbf{x} | C_k) \mathbb{P}(C_k)}{\sum_{j=1}^K \mathbb{P}(\mathbf{x} | C_j) \mathbb{P}(C_j)}$$

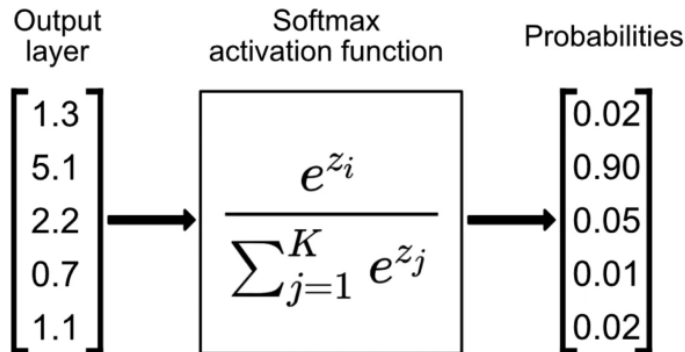
## Multi-class logistic regression (cont.)

- Softmax function **smoothly** highlights the maximum probability and is differentiable.
- Compare it with  $\max(\cdot)$  function which is strict and non-differentiable
- Softmax can also handle negative values because we are using exponential function
- And it gives us probability for each class since:

$$\sum_{k=1}^K \frac{\exp(\mathbf{w}_k^\top \mathbf{x})}{\sum_{j=1}^K \exp(\mathbf{w}_j^\top \mathbf{x})} = 1$$

## Multi-class logistic regression (cont.)

- An example of applying softmax (note that  $z_i = \mathbf{w}^\top \mathbf{x}^{(i)}$ ):



# Multi-class logistic regression (cont.)

- Again we set  $J(\mathbf{W})$  as negative of log likelihood.
- We need  $\mathbf{W}_{\text{MLE}} = \underset{\mathbf{W}}{\text{argmin}} J(\mathbf{W})$

$$\begin{aligned}
 J(W) &= -\log \prod_{i=1}^n \mathbb{P}(\mathbf{y}^{(i)} | \mathbf{x}^{(i)}, \mathbf{W}) \\
 &= -\log \prod_{i=1}^n \prod_{k=1}^K \sigma_k(\mathbf{x}^{(i)}; \mathbf{W})^{y_k^{(i)}} \\
 &= -\sum_{i=1}^n \sum_{k=1}^K y_k^{(i)} \log(\sigma_k(\mathbf{x}^{(i)}; \mathbf{W}))
 \end{aligned}$$

- If **i-th** sample belongs to class  $k$  then  $y_k^{(i)}$  is 1 else 0.
- Again no closed-form solution for  $\mathbf{W}_{\text{MLE}}$

# Multi-class logistic regression (cont.)

- From previous slides we have:

$$J(\mathbf{W}) = - \sum_{i=1}^n \sum_{k=1}^K y_k^{(i)} \log(\sigma_k(\mathbf{x}^{(i)}; \mathbf{W}))$$

- In which:

$$\mathbf{W} = [\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_K], \quad \mathbf{Y} = \begin{pmatrix} \mathbf{y}^{(1)} \\ \mathbf{y}^{(2)} \\ \vdots \\ \mathbf{y}^{(n)} \end{pmatrix} = \begin{pmatrix} y_1^{(1)} & \dots & y_K^{(1)} \\ y_1^{(2)} & \dots & y_K^{(2)} \\ \vdots & \ddots & \vdots \\ y_1^{(n)} & \dots & y_K^{(n)} \end{pmatrix}$$

- $\mathbf{y}$  is a vector of length  $K$  (1-of- $K$  encoding)
  - For example  $\mathbf{y} = [0, 0, 1, 0]^\top$  when the target class is  $C_3$ .



# Multi-class logistic regression (cont.)

- Update rule for gradient descent:

$$\mathbf{w}_j^{t+1} = \mathbf{w}_j^t - \eta \nabla_{\mathbf{w}} J(\mathbf{W}^t)$$
$$\nabla_{\mathbf{w}_j} J(\mathbf{W}) = \sum_{i=1}^n (\sigma_j(\mathbf{x}^{(i)}; \mathbf{W}) - y_j^{(i)}) \mathbf{x}^{(i)}$$

- $\mathbf{w}_j^t$  denotes the weight vector for class  $j$  (since in multi-class LR, each class has its own weight vector) in the  $t$ -th iteration

- ① Introduction
- ② Decision surface
- ③ MLE and MAP
- ④ Gradient descent
- ⑤ Multi-class logistic regression
- ⑥ References**

# Contributions

- **These slides are authored by:**
  - Alireza Mirrokni
  - Danial Gharib
  - Amir Malek Hosseini
  - Aida Jalali

- [1] M. Soleymani Baghshah, “Machine learning.” Lecture slides.
- [2] A. Ng, “ML-005, lecture 6.” Lecture slides.
- [3] C. M. Bishop, *Pattern Recognition and Machine Learning*.  
Information Science and Statistics, New York, NY: Springer, 1 ed., Aug. 2006.
- [4] S. Fidler, “Csc411.” Lecture slides.
- [5] A. Ng and T. Ma, *CS229 Lecture Notes*.  
Updated June 11, 2023.