

# Machine Learning (CE 40717)

## Fall 2025

Ali Sharifi-Zarchi

CE Department  
Sharif University of Technology

November 10, 2025



## 1 Introduction

## 2 Bagging

## 3 Boosting

## 4 AdaBoost

## 5 Comparison

## 6 References

## 1 Introduction

## Condorcet's jury theorem

## Ensemble learning

## Ensemble Methods

## ② Bagging

## ③ Boosting

4 AdaBoost

## 5 Comparison

## 6 References

## 1 Introduction

## Condorcet's jury theorem

## Ensemble learning

## Ensemble Methods

## ② Bagging

## ③ Boosting

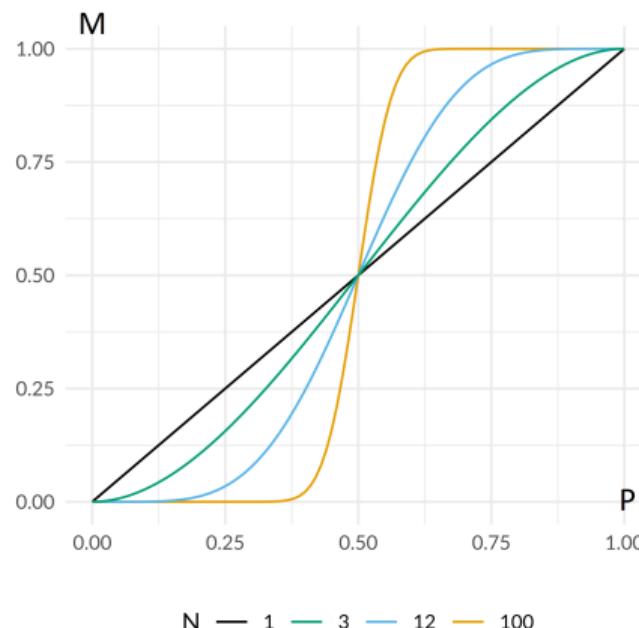
## 4 AdaBoost

## 5 Comparison

## 6 References

## Condorcet's jury theorem

- $N$  voters decide by **majority vote**.
  - $\mathbb{P}(\text{vote correct}) = p$  independently for each voter.
  - Let  $M = \mathbb{P}(\text{majority correct})$ .
  - If  $p > 0.5$  and  $N \rightarrow \infty$ , then  $M \rightarrow 1$ .
    - Intuition: averaging many slightly-better-than-chance votes.



Adapted from Wikipedia

## 1 Introduction

## Condorcet's jury theorem

## Ensemble learning

## Ensemble Methods

## ② Bagging

## ③ Boosting

## 4 AdaBoost

## 5 Comparison

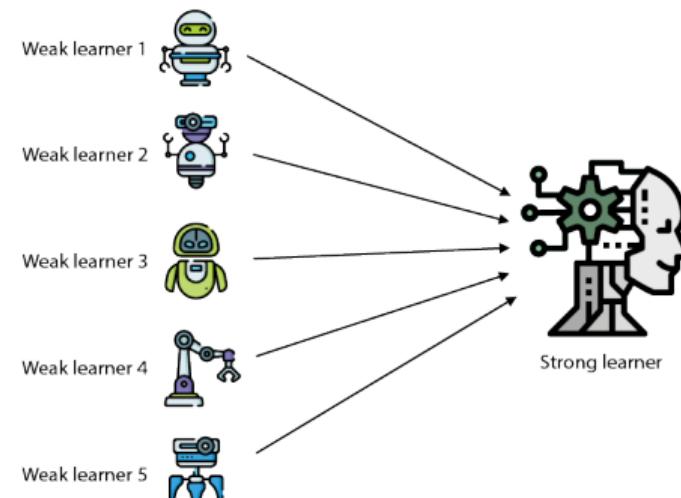
## 6 References

## Strong vs. weak Learners

- **Strong learner:** we seek to produce one classifier for which the classification error can be made arbitrarily small.
    - So far we were looking for such methods.
  - **Weak learner:** a classifier which is just better than random guessing (for now this will be our only expectation).

## Basic idea

- Certain **weak learners** do well in modeling one aspect of the data, while others do well in modeling another.
  - Learn several simple models and **combine** their outputs to produce the final decision.
  - A **composite prediction** where the final accuracy is **better** than the accuracy of **individual models**.



Adapted from [4]

## 1 Introduction

Condorcet's jury theorem

Ensemble learning

Ensemble Methods

## 2 Bagging

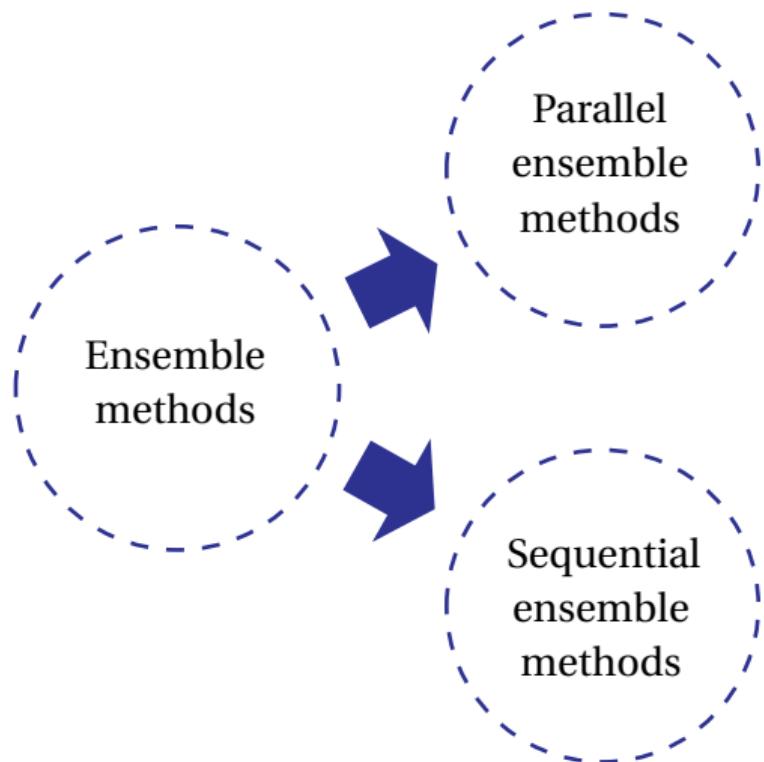
## 3 Boosting

## 4 AdaBoost

## 5 Comparison

## 6 References

# Ensemble Methods



- Weak learners are generated in **parallel**.
- Basic motivation is to use **independence** between the learners.
- Weak learners are generated **consecutively**.
- Basic motivation is to use **dependence** between the base learners.

# What we talk about

- Weak or simple learners
  - **Low variance:** they don't usually overfit
  - **High bias:** they can't learn complex functions
- **Bagging** (parallel): To decrease the variance
  - Random Forest
- **Boosting** (sequential): To decrease the bias (enhance their capabilities)
  - AdaBoost

## 1 Introduction

## 2 Bagging

Basic idea & algorithm

Decision tree (quick review)

Random Forest

## 3 Boosting

## 4 AdaBoost

## 5 Comparison

## 6 References

## 1 Introduction

## 2 Bagging

Basic idea & algorithm

Decision tree (quick review)

Random Forest

## 3 Boosting

## 4 AdaBoost

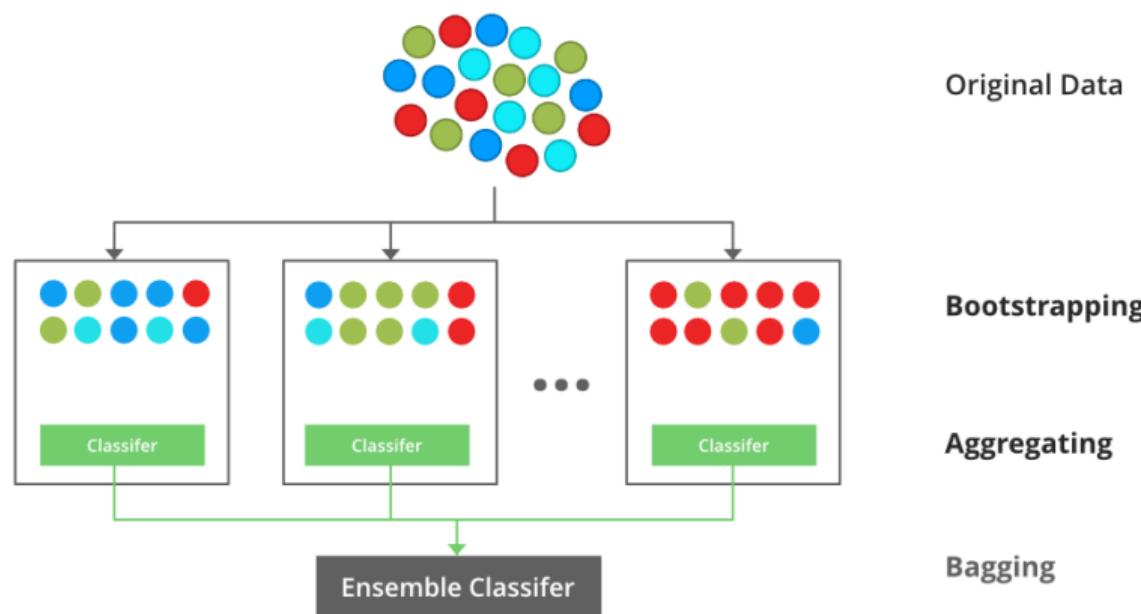
## 5 Comparison

## 6 References

## Basic idea

- **Bagging** = Bootstrap **aggregating**
  - It uses **bootstrap resampling** to generate different training datasets from the original training dataset.
    - Samples training data uniformly at random with replacement.
  - On the training datasets, it trains different weak learners.
  - During testing, it **aggregates** the weak learners by uniform averaging or majority voting.
    - Works best with unstable models (high variance models). Why?

## Basic idea, Cont.



Adapted from GeeksForGeeks

## Algorithm

---

**Algorithm 1** Bagging

- ```

1: Input:  $M$  (required ensemble size),  $D = \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})\}$  (training set)
2: for  $t = 1$  to  $M$  do
3:   Build a dataset  $D_t$  by sampling  $N$  items randomly with replacement from  $D$ 
    ▷ Bootstrap resampling: like rolling an  $N$ -sided die  $N$  time
4:   Train a model  $h_t$  using  $D_t$  and add it to the ensemble
5: end for
6:  $H(\mathbf{x}) = \text{sign}\left(\sum_{t=1}^M h_t(\mathbf{x})\right)$ 
    ▷ Aggregate models by voting for classification or by averaging for regression

```

## 1 Introduction

## 2 Bagging

Basic idea & algorithm

Decision tree (quick review)

Random Forest

## 3 Boosting

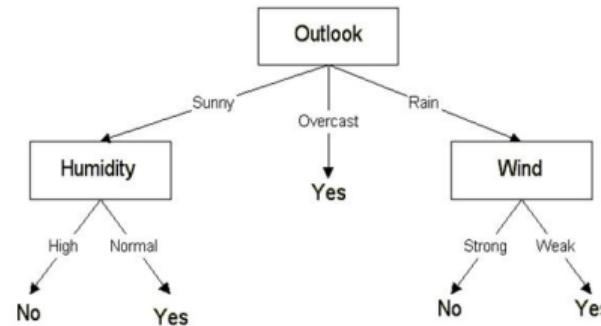
## 4 AdaBoost

## 5 Comparison

## 6 References

# Structure

- **Terminal nodes** (leaves) represent target variable.
- Each **internal node** denotes a test on an attribute.



| Outlook  | Temperature | Humidity | Wind   | Played football(yes/no) |
|----------|-------------|----------|--------|-------------------------|
| Sunny    | Hot         | High     | Weak   | No                      |
| Sunny    | Hot         | High     | Strong | No                      |
| Overcast | Hot         | High     | Weak   | Yes                     |
| Rain     | Mild        | High     | Weak   | Yes                     |
| Rain     | Cool        | Normal   | Weak   | Yes                     |
| Rain     | Cool        | Normal   | Strong | No                      |
| Overcast | Cool        | Normal   | Strong | Yes                     |
| Sunny    | Mild        | High     | Weak   | No                      |
| Sunny    | Cool        | Normal   | Weak   | Yes                     |
| Rain     | Mild        | Normal   | Weak   | Yes                     |
| Sunny    | Mild        | Normal   | Strong | Yes                     |
| Overcast | Mild        | High     | Strong | Yes                     |
| Overcast | Hot         | Normal   | Weak   | Yes                     |
| Rain     | Mild        | High     | Strong | No                      |

Adapted from Medium

# Learning

- Learning an optimal decision tree is **NP-Complete**.
  - Instead, we use a **greedy search** based on a heuristic.
  - We can't guarantee to return the globally-optimal decision tree.
- The most common strategy for DT learning is a greedy top-down approach.
- Tree is constructed by splitting samples into subsets based on an **attribute value test** in a recursive manner.

Adapted from G.E. Naumov, "NP-completeness of problems of construction of optimal decision trees", 1991

# Algorithm

---

## Algorithm 2 Constructing DT

---

```
1: procedure FINDTREE( $S, A$ )                                ▷ Input:  $S$  (samples),  $A$  (attributes)
2:   if  $A$  is empty or all labels in  $S$  are the same then
3:     status  $\leftarrow$  leaf
4:     class  $\leftarrow$  most common class in  $S$ 
5:   else
6:     status  $\leftarrow$  internal
7:      $a \leftarrow \text{bestAttribute}(S, A)$                                 ▷ The attribute value test
8:     LeftNode  $\leftarrow \text{FindTree}(S(a > t), A - \{a\})$                 ▷  $t$  (threshold)
9:     RightNode  $\leftarrow \text{FindTree}(S(a \leq t), A - \{a\})$ 
10:   end if
11: end procedure
```

---

# Which attribute is the best?

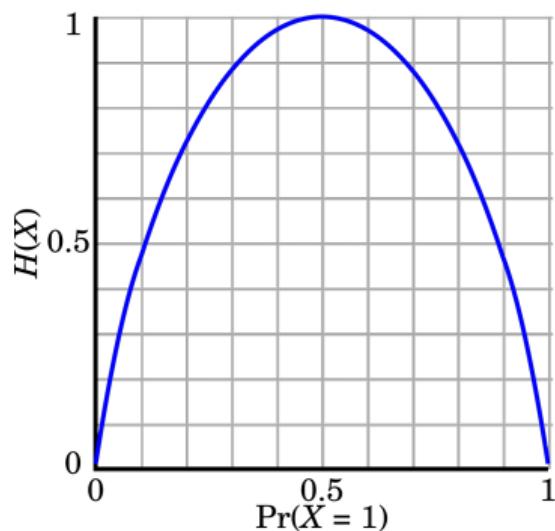
- **Entropy** measures label uncertainty.

$$H_S(Y) = - \sum_{y \in \text{Values}(Y)} \mathbb{P}_S(Y=y) \log_2 \mathbb{P}_S(Y=y)$$

- **Information Gain (IG)**

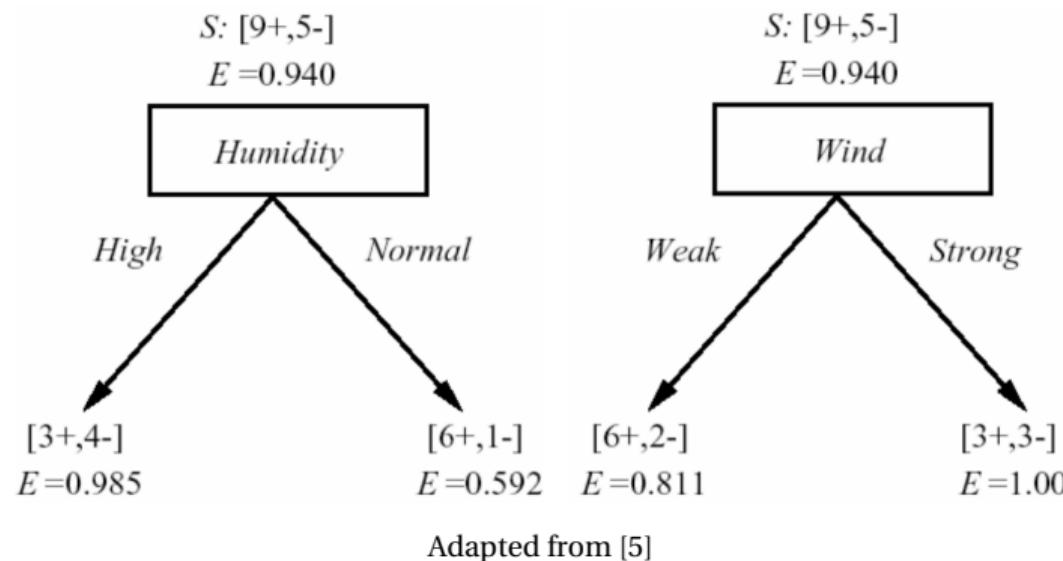
$$\text{Gain}(S, A) = H_S(Y) - \sum_{v \in \text{Values}(A)} \frac{|S_v|}{|S|} H_{S_v}(Y)$$

$A$ : splitting attribute     $Y$ : target     $S$ : samples     $S_v$ : subset with  $A = v$



Adapted from Wikipedia

## Example



$$\text{Gain}(S, \text{Humidity}) = 0.940 - (7/14)0.985 - (7/14)0.592 = 0.151$$

$$\text{Gain}(S, \text{Wind}) = 0.940 - (8/14) 0.811 - (6/14) 1.0 = 0.048$$

## 1 Introduction

## 2 Bagging

Basic idea & algorithm

Decision tree (quick review)

Random Forest

## 3 Boosting

## 4 AdaBoost

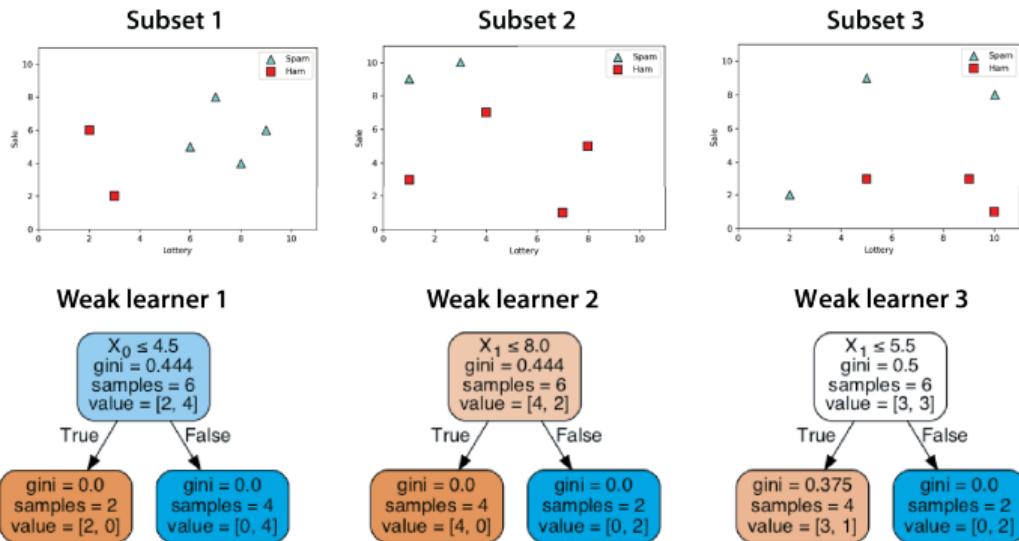
## 5 Comparison

## 6 References

## Bagging on decision trees?

## Why decision trees?

- Interpretable
  - Robust to outliers
  - **Low bias**
  - **High variance**



Adapted from [4]

# Perfect candidates

- Why are **DTs** good candidates for ensembles?
  - Consider averaging many (nearly) **unbiased** tree estimators.
  - Bias remains similar, but **variance is reduced**.
- Remember Bagging?
  - Train many trees on bootstrapped data, then aggregate (average/majority) the outputs.

# Algorithm

---

## Algorithm 3 Random Forest

---

- 1: **Input:**  $T$  (number of trees),  $m$  (number of variables used to split each node)
  - 2: **for**  $t = 1$  to  $T$  **do**
  - 3:     Draw a bootstrap dataset
  - 4:     At each node, sample  $m$  features uniformly from  $\{1, \dots, d\}$  as split candidates
  - 5: **end for**
  - 6: **Output:** ▷ Usually:  $m \leq \sqrt{d}$
  - 7:     Regression: average of the outputs
  - 8:     Classification: majority voting
-

## Why Random Forest Works: Variance & Correlation

Consider  $M$  base learners with variance  $\sigma^2$  and pairwise correlation  $\rho$ . For the averaged predictor  $\bar{h}(\mathbf{x}) = \frac{1}{M} \sum_{t=1}^M h_t(\mathbf{x})$ :

$$\text{Var}(\bar{h}) = \rho \sigma^2 + \frac{1 - \rho}{M} \sigma^2$$

- **Bagging** reduces  $\frac{1-\rho}{M}\sigma^2$  via averaging ( $M \uparrow$ ).
- **Random Forest** also lowers  $\rho$  by random feature subspacing at each split.
- Takeaway: smaller  $\rho$  and larger  $M \Rightarrow$  lower ensemble variance without increasing bias much.

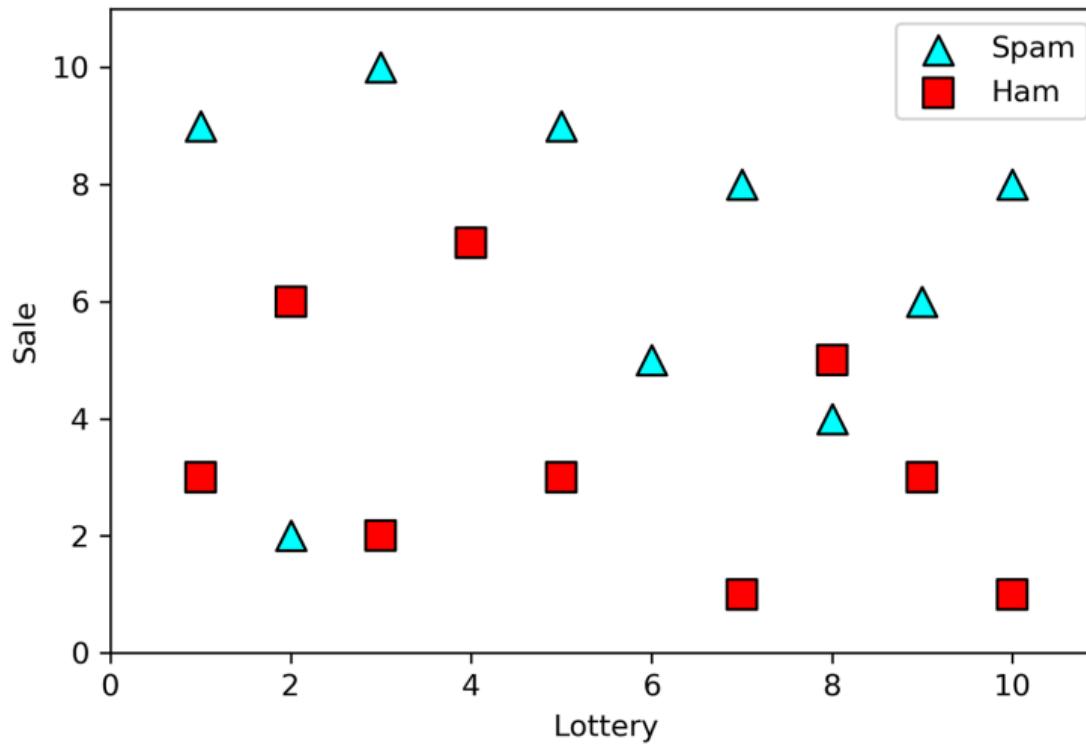
# Out-of-Bag (OOB) Estimation

- In each bootstrap sample of size  $N$ , the probability an example is *not* selected is  $(1 - \frac{1}{N})^N \approx e^{-1} \approx 36.8\%$ .
- Use OOB samples for each tree to estimate generalization error *without* a separate validation split.
- Practical tips:
  - Plot OOB error vs. number of trees; stop once the curve flattens.
  - OOB can also support permutation-based feature importance.

# Feature Importance in Forests (MDI vs Permutation)

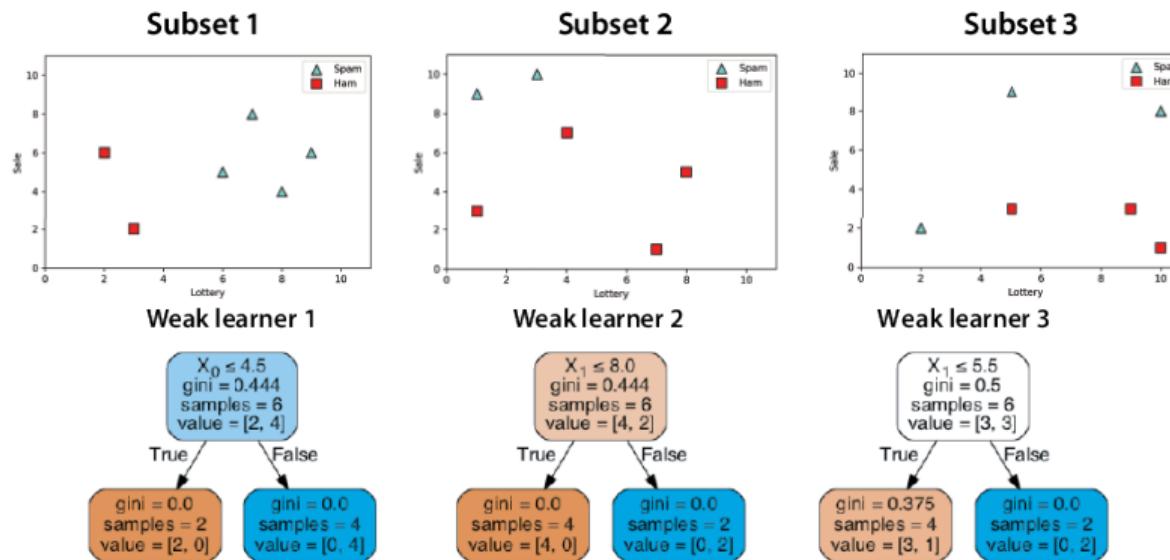
- **MDI (Gini/Entropy decrease)**: fast, but biased toward high-cardinality or noisy features.
- **Permutation importance (OOB)**: shuffle a feature and measure OOB error increase.
- Guidelines:
  - Prefer permutation importance for **comparisons** across features.
  - Beware correlation: grouped or conditional permutation can help.
  - For deeper interpretability, use PDP/ICE; for local attributions, consider SHAP (optional).

## Example



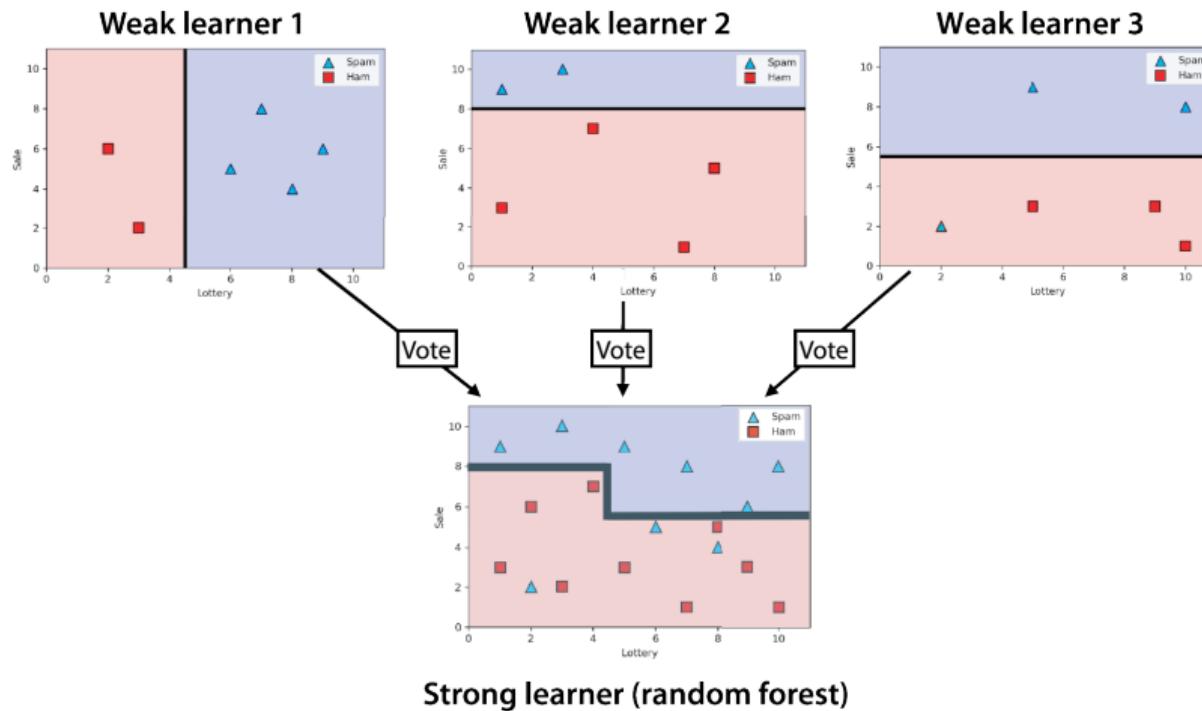
Adapted from [4]

## Example, Cont.



Adapted from [4]

## Example, Cont.



Adapted from [4]

## 1 Introduction

## 2 Bagging

## 3 Boosting

Motivation & basic idea

Algorithm

## 4 AdaBoost

## 5 Comparison

## 6 References

## 1 Introduction

## 2 Bagging

## 3 Boosting

Motivation & basic idea

Algorithm

## 4 AdaBoost

## 5 Comparison

## 6 References

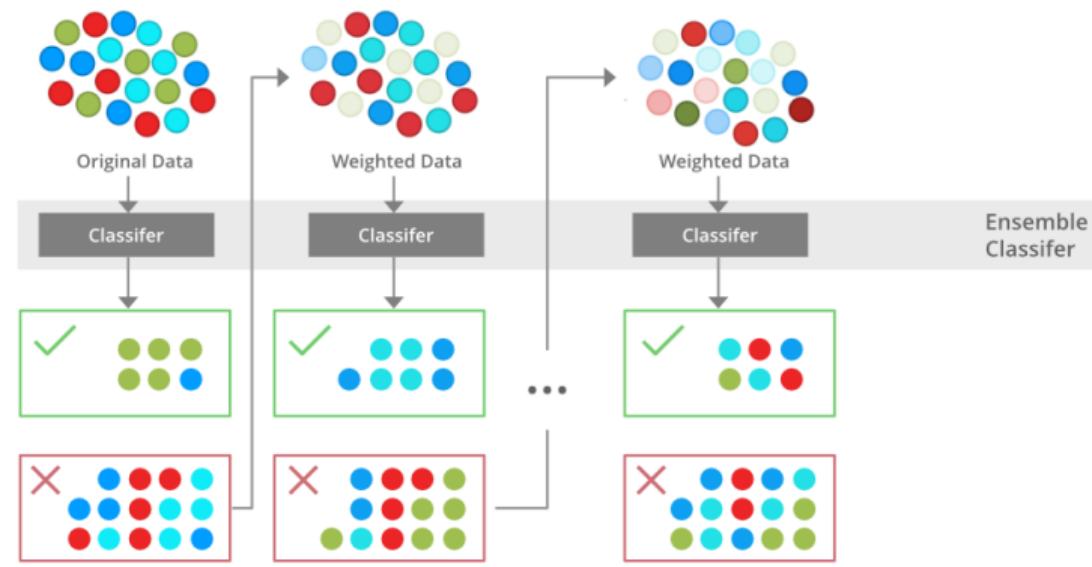
# Problems with bagging

- Bagging created a diversity of **weak learners** by creating random datasets.
  - Examples: Decision stumps (shallow decision trees), Logistic regression, ...
- Did we have full control over the usefulness of the weak learners?
  - The **diversity** or **complementarity** of the weak learners is not controlled in any way, it is left to chance and to the instability (variance) of the models.

## Basic idea

- We would expect a better performance if the weak learners also **complemented** each other.
  - They would have “expertise” on different subsets of the dataset.
  - So they would work better on different subsets.
- The basic idea of boosting is to generate a **series** of weak learners which complement each other.
  - For this, we will force each learner to focus on **the mistakes of the previous learner**.

## Basic idea, Cont.



Adapted from GeeksForGeeks

## 1 Introduction

## 2 Bagging

## 3 Boosting

Motivation & basic idea

Algorithm

## 4 AdaBoost

## 5 Comparison

## 6 References

## Algorithm

- Try to combine many simple **weak** learners (in sequence) to find a single **strong** learner (For simplicity, suppose that we have a classification problem from now on).
    - Each component is a simple binary  $\pm 1$  classifier
    - Voted combinations of component classifiers

$$H_M(\mathbf{x}) \equiv \alpha_1 h(\mathbf{x}; \boldsymbol{\theta}_1) + \cdots + \alpha_M h(\mathbf{x}; \boldsymbol{\theta}_M)$$

- To simplify notations:  $h(\mathbf{x}; \boldsymbol{\theta}_i) = h_i(\mathbf{x})$

$$H_M(\mathbf{x}) = \alpha_1 h_1(\mathbf{x}) + \cdots + \alpha_M h_M(\mathbf{x})$$

- **Prediction:**  $\hat{y} = \text{sign}(H_M(\mathbf{x}))$

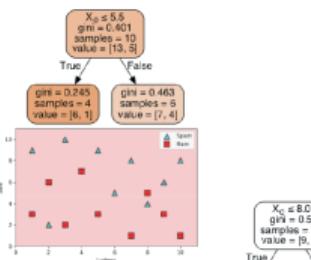
# Candidate for $h_i(\mathbf{x})$

- **Decision stumps**
- Each classifier is based on only a single feature of  $\mathbf{x}$  (e.g.,  $x_k$ ):

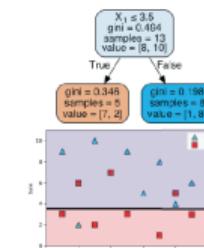
$$h(\mathbf{x}; \boldsymbol{\theta}) = \text{sign}(w_1 x_k - w_0)$$

$$\boldsymbol{\theta} = \{k, w_1, w_0\}$$

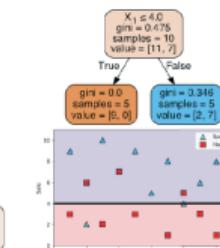
**Weak learner 1**



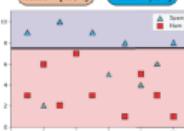
**Weak learner 3**



**Weak learner 5**



**Weak learner 2**



Adapted from [4]

## Gradient Boosting (Functional Gradient Descent)

- View boosting as minimizing a loss  $L(y, F(\mathbf{x}))$  by **steepest descent in function space**.

$$F_m(\mathbf{x}) = F_{m-1}(\mathbf{x}) + v \beta_m T_m(\mathbf{x}), \quad r_{im} = -\left[ \frac{\partial L(y_i, F)}{\partial F} \right]_{F=F_{m-1}(\mathbf{x}^{(i)})}$$

- Fit a tree  $T_m$  to pseudo-residuals  $\{r_{im}\}$ , choose  $\beta_m$  by line search;  $v \in (0, 1]$  is the **learning rate**.
  - Choice of  $L$ : square loss (regression), logistic loss (classification), etc.
  - Regularization: small  $v$ , shallow trees, subsampling of rows/columns.

## 1 Introduction

## 2 Bagging

## 3 Boosting

## 4 AdaBoost

Basic idea & algorithm

Loss function & proof

Properties (extra-reading)

## 5 Comparison

## 6 References

## 1 Introduction

## 2 Bagging

## 3 Boosting

## 4 AdaBoost

Basic idea & algorithm

Loss function & proof

Properties (extra-reading)

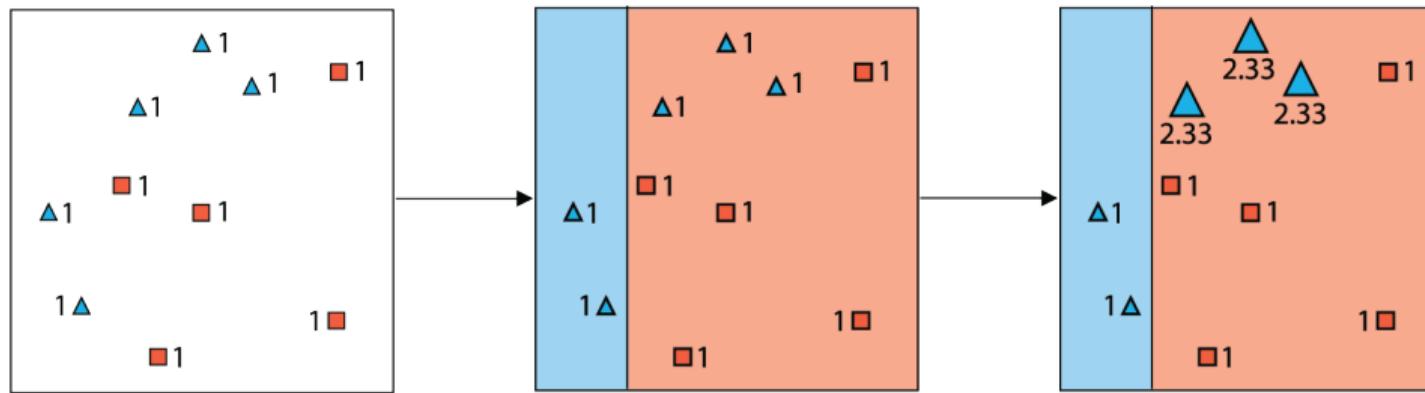
## 5 Comparison

## 6 References

## Basic idea

- **Sequential** production of classifiers
  - Iteratively add the classifier whose addition will be most helpful.
- Represent the importance of each sample by assigning **weights** to them.
  - Correct classification  $\Rightarrow$  smaller weights
  - Misclassified samples  $\Rightarrow$  larger weights
- Each classifier is **dependent** on the previous ones.
  - Focuses on the **previous ones' error**.

## Example



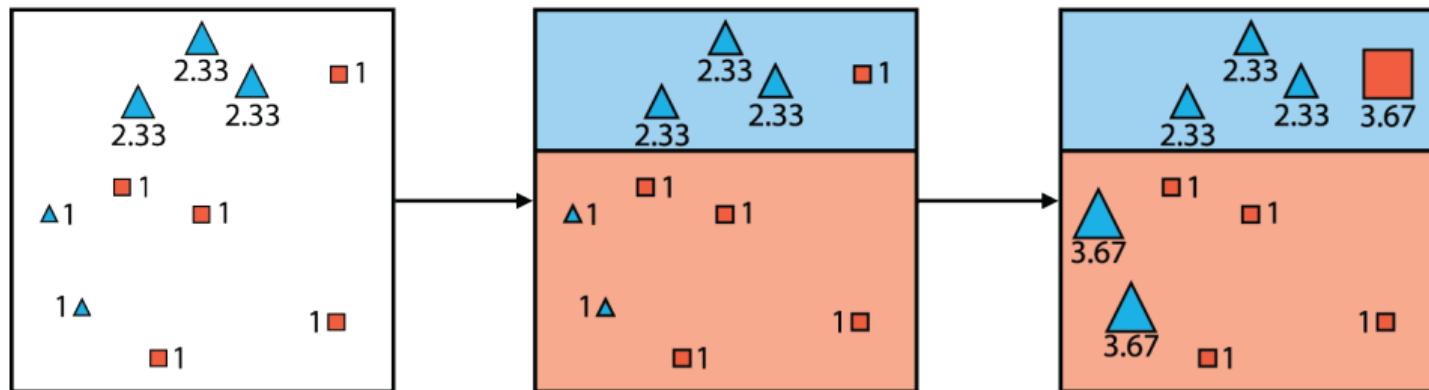
Add a weight of 1  
to every point.

Fit a weak learner.  
Correct: 7  
Incorrect: 3

Adapted from [4]

Rescale misclassified  
points by 7/3.

## Example, Cont.



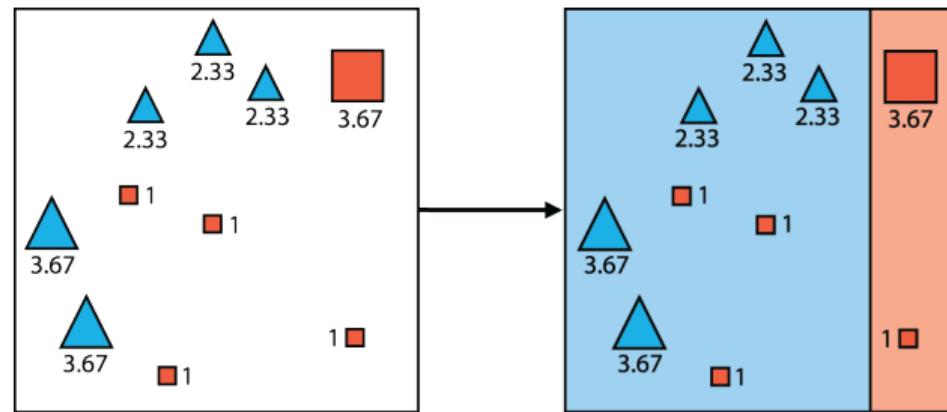
The rescaled dataset

Fit a weak learner.  
Sum of correct: 11  
Sum of incorrect: 3

Rescale misclassified points by 11/3.

Adapted from [4]

## Example, Cont.

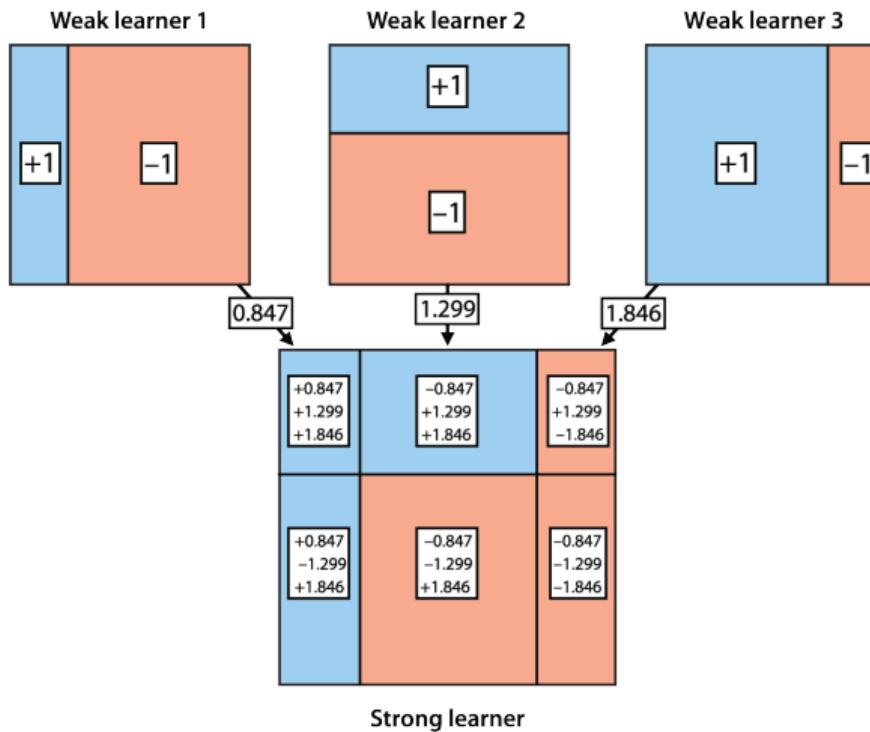


## The rescaled dataset

Fit a weak learner.  
Sum of correct: 19  
Sum of incorrect: 3

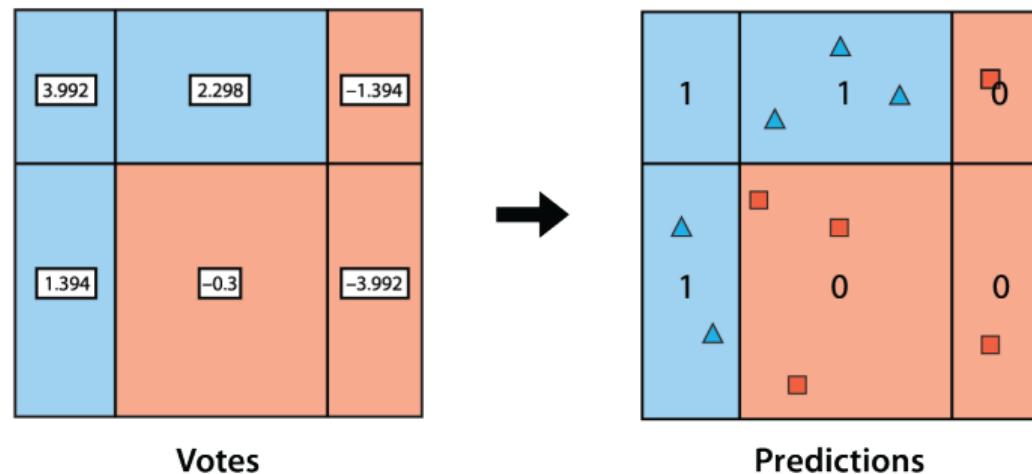
Adapted from [4]

## Example, Cont.



Adapted from [4]

## Example, Cont.



Adapted from [4]

## Algorithm

- $H_M(\mathbf{x}) = \frac{1}{2}[\alpha_1 h_1(\mathbf{x}) + \cdots + \alpha_M h_M(\mathbf{x})] \longrightarrow$  the complete model  $y^{(i)} \in \{-1, 1\}$
  - $h_m(\mathbf{x})$ :  $m$ -th weak learner
  - $\alpha_m = ? \longrightarrow$  votes of the  $m$ -th weak learner
  - $w_m^{(i)}$ : weight of sample  $i$  in iteration  $m$ 
    - $w_{m+1}^{(i)} = ?$
  - $J_m = \sum_{i=1}^N w_m^{(i)} \times \mathbb{1}(y^{(i)} \neq h_m(\mathbf{x}^{(i)})) \longrightarrow$  loss of the  $m$ -th weak learner
  - $\epsilon_m = \frac{\sum_{i=1}^N w_m^{(i)} \times \mathbb{1}(y^{(i)} \neq h_m(\mathbf{x}^{(i)}))}{\sum_{i=1}^N w_m^{(i)}} \longrightarrow$  weighted error of the  $m$ -th weak learner

## Algorithm, Cont.

- $H_M(\mathbf{x}) = \frac{1}{2}[\alpha_1 h_1(\mathbf{x}) + \cdots + \alpha_M h_M(\mathbf{x})] \longrightarrow$  the complete model  $y^{(i)} \in \{-1, 1\}$
  - $h_m(\mathbf{x})$ :  $m$ -th weak learner
  - $\alpha_m = \ln\left(\frac{1 - \epsilon_m}{\epsilon_m}\right)$   $\longrightarrow$  votes of the  $m$ -th weak learner
  - $w_m^{(i)}$ : weight of sample  $i$  in iteration  $m$ 
    - $w_{m+1}^{(i)} = w_m^{(i)} e^{\alpha_m \mathbb{1}(y^{(i)} \neq h_m(\mathbf{x}^{(i)}))}$
  - $J_m = \sum_{i=1}^N w_m^{(i)} \times \mathbb{1}(y^{(i)} \neq h_m(\mathbf{x}^{(i)})) \longrightarrow$  loss of the  $m$ -th weak learner
  - $\epsilon_m = \frac{\sum_{i=1}^N w_m^{(i)} \times \mathbb{1}(y^{(i)} \neq h_m(\mathbf{x}^{(i)}))}{\sum_{i=1}^N w_m^{(i)}} \longrightarrow$  weighted error of the  $m$ -th weak learner

## Algorithm, Cont.

---

**Algorithm 4** AdaBoost

- 1: Initialize data weight  $w_1^{(i)} = \frac{1}{N}$  for all  $N$  samples

2: for  $m = 1$  to  $M$  do

- 3: Find  $h_m(\mathbf{x})$  by minimizing the loss:

$$J_m = \sum_{i=1}^N w_m^{(i)} \times \mathbb{1}(y^{(i)} \neq h_m(\mathbf{x}^{(i)}))$$

- 4: Find the weighted error of  $h_m(\mathbf{x})$ .

$$\epsilon_m = \frac{\sum_{i=1}^N w_m^{(i)} \times \mathbb{1}(y^{(i)} \neq h_m(\mathbf{x}^{(i)}))}{\sum_{i=1}^N w_m^{(i)}}$$

- 5: Assign votes  $\alpha_m = \ln\left(\frac{1 - \epsilon_m}{\epsilon_m}\right)$

$$w_{m+1}^{(i)} = w_m^{(i)} e^{\alpha_m \mathbb{1}_{\{y^{(i)} \neq h_m(\mathbf{x}^{(i)})\}}}$$

- 6: Update the weights:

7: end for

- 8: **Combined classifier:**  $\hat{y} = \text{sign}(H_M(\mathbf{x}))$  where  $H_M(\mathbf{x}) = \frac{1}{2} \sum_{m=1}^M \alpha_m h_m(\mathbf{x})$

## 1 Introduction

## 2 Bagging

## 3 Boosting

## 4 AdaBoost

Basic idea & algorithm

Loss function & proof

Properties (extra-reading)

## 5 Comparison

## 6 References

## Loss function

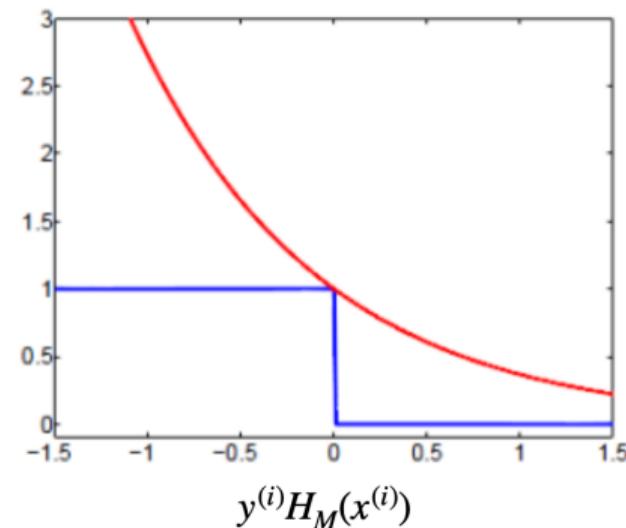
- There are many options for the loss function.
    - AdaBoost is equivalent to using the following **exponential loss**.

$$\mathcal{L}(y, H_M(\mathbf{x})) = e^{-y \times H_M(\mathbf{x})}$$

$$\hat{y} = \text{sign}(H_M(\mathbf{x}))$$

# Why the exponential loss?

- Differentiable approximation (bound) of the **0/1 loss**
  - Easy to optimize
  - Optimizing an upper bound on classification error.



Adapted from [2]

### Step 1: Calculating the exponential loss

- We need to calculate the exponential loss for

$$H_m(\mathbf{x}) = \frac{1}{2}[\alpha_1 h_1(\mathbf{x}) + \cdots + \alpha_m h_m(\mathbf{x})]$$

To have a cleaner form later

- **Idea:** consider adding the  $m$ -th component

$$\begin{aligned} \mathcal{L}_m &= \sum_{i=1}^N e^{-y^{(i)} H_m(\mathbf{x}^{(i)})} = \sum_{i=1}^N e^{-y^{(i)} [H_{m-1}(\mathbf{x}^{(i)}) + \frac{1}{2} \alpha_m h_m(\mathbf{x}^{(i)})]} \\ &= \sum_{i=1}^N e^{-y^{(i)} H_{m-1}(\mathbf{x}^{(i)})} e^{-\frac{1}{2} \alpha_m y^{(i)} h_m(\mathbf{x}^{(i)})} = \sum_{i=1}^N \underbrace{w_m^{(i)}}_{e^{-y^{(i)} H_{m-1}(\mathbf{x}^{(i)})}} e^{-\frac{1}{2} \alpha_m y^{(i)} h_m(\mathbf{x}^{(i)})} \end{aligned}$$

Suppose it is fixed at stage  $m$

Should be optimized at stage  $m$  by seeking  $h_m(\mathbf{x})$  and  $\alpha_m$

## Step 2: Deriving the weighted error function

- We need to derive the weighted error function,  $J_m$

$$\begin{aligned}
\mathcal{L}_m &= \sum_{i=1}^N w_m^{(i)} e^{-\frac{1}{2}\alpha_m y^{(i)} h_m(\mathbf{x}^{(i)})} \\
&= e^{\frac{-\alpha_m}{2}} \left( \sum_{y^{(i)} = h_m(\mathbf{x}^{(i)})} w_m^{(i)} \right) + e^{\frac{\alpha_m}{2}} \left( \sum_{y^{(i)} \neq h_m(\mathbf{x}^{(i)})} w_m^{(i)} \right) \\
&= (e^{\frac{\alpha_m}{2}} - e^{\frac{-\alpha_m}{2}}) \underbrace{\left( \sum_{y^{(i)} \neq h_m(\mathbf{x}^{(i)})} w_m^{(i)} \right)}_{J_m} + e^{\frac{-\alpha_m}{2}} \left( \sum_{i=1}^N w_m^{(i)} \right)
\end{aligned}$$

Find  $h_m(\mathbf{x})$  that minimizes  $J_m$

## Step 3: Deriving $\epsilon_m$ and $\alpha_m$

- We need to derive  $\epsilon_m$  and  $\alpha_m$  by setting the derivative equal to zero:

$$\frac{\partial \mathcal{L}_m}{\partial \alpha_m} = 0$$

- Idea:** separate the derivative into misclassified and correctly classified samples.

$$\begin{aligned}\implies & \frac{1}{2}(e^{\frac{\alpha_m}{2}} + e^{-\frac{\alpha_m}{2}}) \left( \sum_{y^{(i)} \neq h_m(\mathbf{x}^{(i)})} w_m^{(i)} \right) = \frac{1}{2} e^{-\frac{\alpha_m}{2}} \left( \sum_{i=1}^N w_m^{(i)} \right) \\ \implies & \frac{e^{-\frac{\alpha_m}{2}}}{(e^{\frac{\alpha_m}{2}} + e^{-\frac{\alpha_m}{2}})} = \frac{\sum_{y^{(i)} \neq h_m(\mathbf{x}^{(i)})} w_m^{(i)}}{\sum_{i=1}^N w_m^{(i)}}\end{aligned}$$

- Set  $\epsilon_m = \frac{\sum_{i=1}^N w_m^{(i)} \mathbb{1}(y^{(i)} \neq h_m(\mathbf{x}^{(i)}))}{\sum_{i=1}^N w_m^{(i)}} \implies \alpha_m = \ln\left(\frac{1-\epsilon_m}{\epsilon_m}\right)$

#### Step 4: Justifying the weight update mechanism

- We need to justify the weight update mechanism.
  - **Idea:** we have  $w_m^{(i)}$  from the first step as  $w_{m+1}^{(i)} = e^{-y^{(i)} H_m(\mathbf{x}^{(i)})}$

$$\xrightarrow{\text{separate } h_m(\mathbf{x}^{(i)})} w_{m+1}^{(i)} = w_m^{(i)} e^{-\frac{1}{2} \alpha_m y^{(i)} h_m(\mathbf{x}^{(i)})}$$

$$\xrightarrow{y^{(i)} h_m(\mathbf{x}^{(i)}) = 1 - 2 \mathbb{I}(y^{(i)} \neq h_m(\mathbf{x}^{(i)}))} w_{m+1}^{(i)} = w_m^{(i)} e^{-\frac{1}{2} \alpha_m} e^{\alpha_m \mathbb{I}(y^{(i)} \neq h_m(\mathbf{x}^{(i)}))}$$

Independent of  $i$  and can be ignored

$$\implies w_{m+1}^{(i)} = w_m^{(i)} e^{\alpha_m \mathbb{1}(y^{(i)} \neq h_m(\mathbf{x}^{(i)}))}$$

## 1 Introduction

## 2 Bagging

## 3 Boosting

## 4 AdaBoost

Basic idea & algorithm

Loss function & proof

Properties (extra-reading)

## 5 Comparison

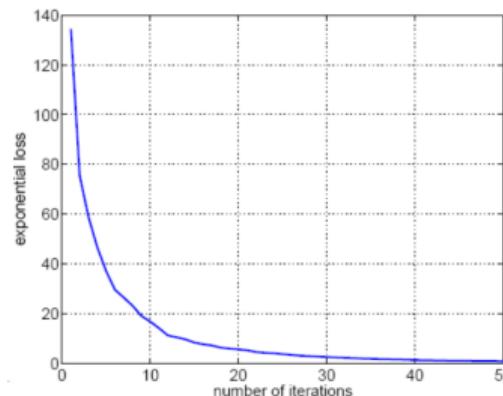
## 6 References

## Exponential loss properties

- In each boosting iteration, assuming we can find  $h(\mathbf{x}; \boldsymbol{\theta}_m)$  whose weighted error is better than chance.

$$H_m(\mathbf{x}) = \frac{1}{2}[\alpha_1 h(\mathbf{x}; \boldsymbol{\theta}_1) + \cdots + \alpha_m h(\mathbf{x}; \boldsymbol{\theta}_m)]$$

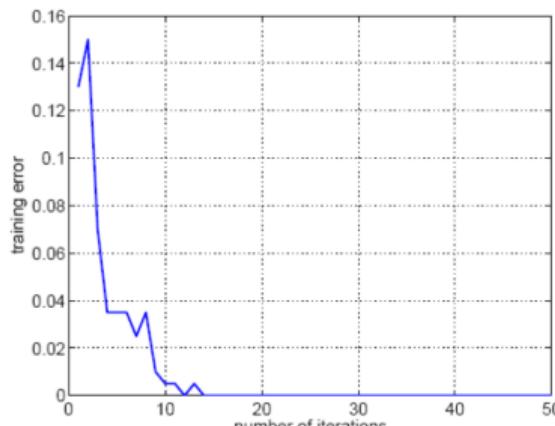
- Thus, **lower exponential loss** over training data is guaranteed.



Adapted from [6]

## Training error properties

- Boosting iterations typically **decrease** the **training error** of  $H_M(\mathbf{x})$  over training examples.

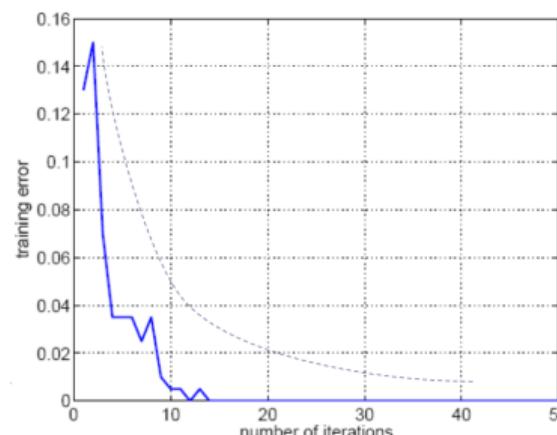


Adapted from [6]

## Training error properties, Cont.

- **Training error** has to go **down exponentially fast** if the weighted error of each  $h_m$  is strictly better than chance (i.e.,  $\epsilon_m < 0.5$ )

$$E_{\text{train}}(H_M) \leq \prod_{m=1}^M 2\sqrt{\epsilon_m(1-\epsilon_m)}$$

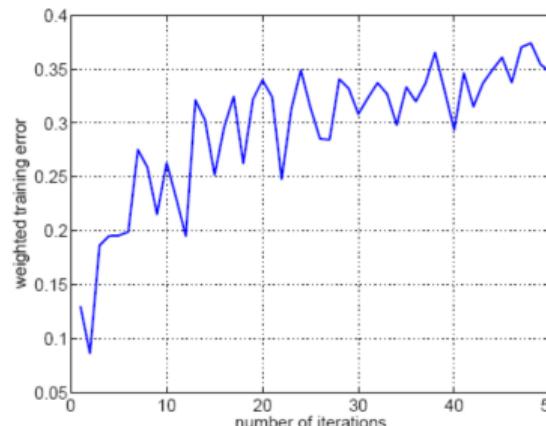


Adapted from [6]

## Weighted error properties

- **Weighted error** of each new component classifier tends to **increase** as a function of boosting iterations.

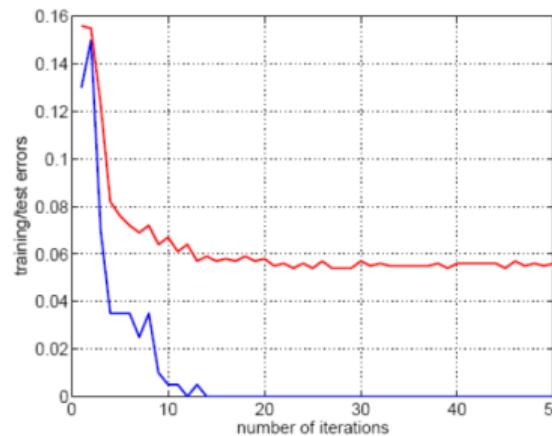
$$\epsilon_m = \frac{\sum_{i=1}^N w_m^{(i)} \mathbb{1}(y^{(i)} \neq h_m(\mathbf{x}^{(i)}))}{\sum_{i=1}^N w_m^{(i)}}$$



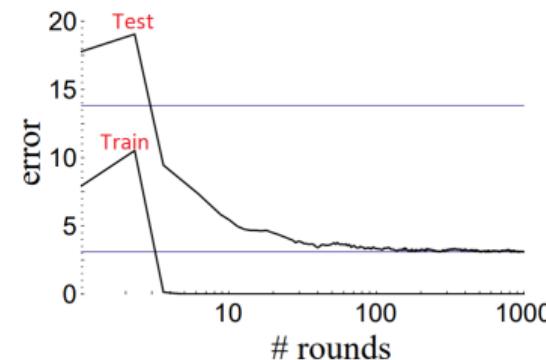
Adapted from [6]

## Test error properties

- **Test error** can still **decrease** after training error is flat (even zero).
  - But, is it robust to overfitting?
    - May easily overfit in the presence of labeling noise or overlap of classes.



Adapted from [6]



Adapted from [3]

# Boosting in Practice: Avoiding Overfitting

- **Label noise & overlap:** exponential loss can overweight hard/noisy points.
- Remedies:
  - Use smoother losses (e.g., logistic) or early stopping on a validation/OOB proxy.
  - Small learning rate  $\nu$  (e.g., 0.05–0.1) with more iterations.
  - Shallow base trees (stumps to depth 3–5) to control variance.
  - Row/feature subsampling to decorrelate learners.
  - Handle class imbalance via class weights or balanced sampling; evaluate with PR-AUC.

## Typical behavior

- **Exponential loss goes strictly down.**
- **Training error of  $H$  goes down.**
- Weighted error  $\epsilon_m$  goes **up**  $\Rightarrow$  share of votes  $\alpha_m$  goes **down**.
- **Test error decreases** even after a flat training error.

## 1 Introduction

## 2 Bagging

## 3 Boosting

## 4 AdaBoost

## 5 Comparison

## 6 References

# Bagging vs. Boosting

|                            | Bagging                             | Boosting                                     |
|----------------------------|-------------------------------------|----------------------------------------------|
| <b>Training Strategy</b>   | Parallel training                   | Sequential training                          |
| <b>Data Sampling</b>       | Bootstrapping<br>(random subsets)   | Weighted<br>(by instance importance)         |
| <b>Learners Dependency</b> | Independent                         | Dependent<br>(on the previous models)        |
| <b>Learner Weighting</b>   | Equal weights                       | Varying weights<br>(based on importance)     |
| <b>Tolerance to Noise</b>  | More robust<br>(due to aggregation) | More sensitive<br>(may overfit to noise)     |
| <b>Properties</b>          | Reduces variance                    | Reduces bias and variance<br>(focus on bias) |

## 1 Introduction

## 2 Bagging

## 3 Boosting

## 4 AdaBoost

## 5 Comparison

## 6 References

# Contributions

- **This slide has been prepared thanks to:**
  - Nikan Vasei
  - Mahan Bayhaghi
  - Aida Jalali

- [1] C. M., *Pattern Recognition and Machine Learning*.  
Information Science and Statistics, New York, NY: Springer, 1 ed., Aug. 2006.
- [2] M. Soleymani Baghshah, “Machine learning.” Lecture slides.
- [3] R. E. Schapire, “The boosting approach to machine learning: An overview,”  
*Nonlinear estimation and classification*, pp. 149–171, 2003.
- [4] L. Serrano, *Grokking machine learning*.  
New York, NY: Manning Publications, Jan. 2022.
- [5] T. Mitchell, *Machine Learning*.  
McGraw-Hill series in computer science, New York, NY: McGraw-Hill Professional,  
Mar. 1997.
- [6] T. Jaakkola, “Machine learning course slides.” Lecture slides.