

# Machine Learning (CE 40717)

Fall 2024

Ali Sharifi-Zarchi

CE Department  
Sharif University of Technology

December 17, 2025



## 1 Optimization

## 2 The Loss Surface

## 3 Gradient Descent

## 4 Momentum

## 5 Newton's optimization Method

## 6 References

## 1 Optimization

## ② The Loss Surface

### 3 Gradient Descent

## 4 Momentum

## 5 Newton's optimization Method

## 6 References

## Optimization Problem

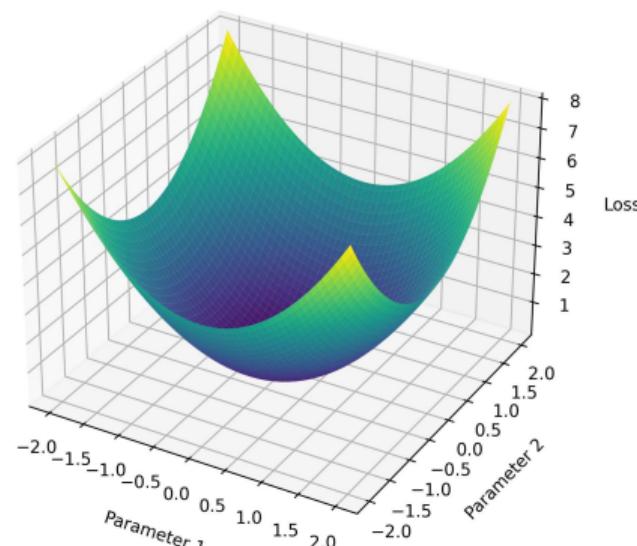
- **Goal:** Find the value of  $x$  where  $f(x)$  is at a **minimum or maximum**.
  - In neural networks, we aim to minimize **prediction error** by finding the optimal weights  $w^*$ :

$$w^* = \arg \min_w J(w)$$

- Simply put: determine the **direction to step** that will quickly **reduce loss**

## Convexity and Optimization

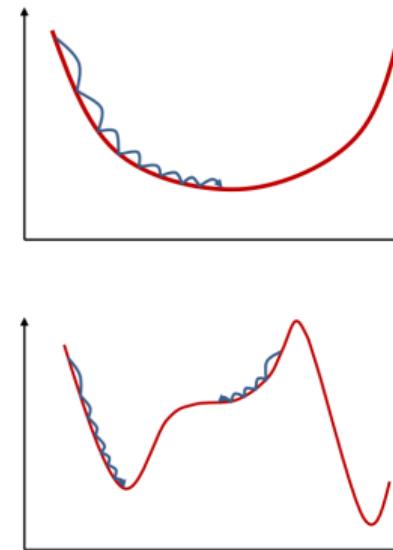
- **Convex Functions:**
    - A function is **convex** if any line segment between points on the curve lies **above or on** the curve.
    - Convex functions are easier to optimize, as they have a single **global minimum**.
    - Numerical methods like **Gradient Descent** are guaranteed to reach the global minimum in convex functions.



**Figure 1:** Example of convex function (bowl shape)

## Non-Convex Functions and Challenges

- **Non-Convex Functions:**
    - Characterized by multiple **local minima** and **saddle points**.
    - **Global Minimum:** Overall lowest point.
    - **Local Minimum:** Lower than nearby points, but not the lowest overall.
    - **Saddle Points:** Regions where the gradient is close to zero but can increase or decrease in other directions.
  - Finding the **global minimum** is more complex in non-convex functions.



**Figure 2:** Convex (top) vs Non-Convex (bottom) functions. Source: (CMU, 11-785)

## 1 Optimization

## 2 The Loss Surface

### 3 Gradient Descent

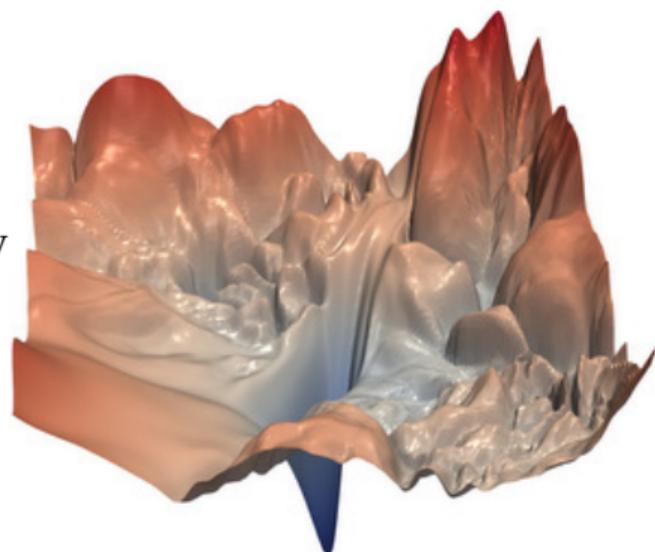
## 4 Momentum

## 5 Newton's optimization Method

## 6 References

## Loss Surface Definition

- The **loss surface** shows how error changes based on network weights.
  - For neural networks, the loss surface is typically **non-convex** due to multiple layers, nonlinear activations, and complex parameter interactions, resulting in **multiple local minima** and **saddle points**.
  - In large networks, most local minima yield similar error values close to the **global minimum**; this is less true in smaller networks.



**Figure 3:** Loss surface of ResNet56. Source: GitHub: Loss Landscape

## Loss Optimization

- **Goal:** How can we optimize a non-convex loss function effectively?
  - **Gradient Descent:**
    - This method identifies the **steepest descent direction** to guide the optimization process.
  - **Newton's Method:**
    - This method looks for **critical points** where the derivative  $f'(x) = 0$ , which may indicate minima, maxima, or saddle points.
    - Newton's Method uses the second derivative (Hessian) to adjust step sizes, which can lead to faster convergence compared to Gradient Descent.

## 1 Optimization

## 2 The Loss Surface

## 3 Gradient Descent

## 4 Momentum

## 5 Newton's optimization Method

## 6 References

## Gradient Descent Overview

- **Gradient Descent:** As mentioned earlier in this course, Gradient Descent is an iterative method to minimize error by updating weights in the direction of the **negative gradient**:

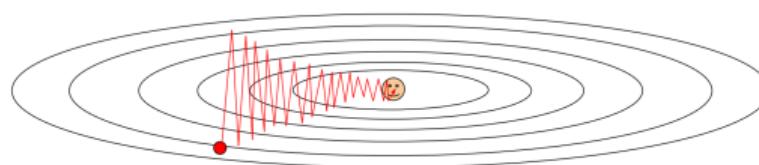
$$w_{t+1} = w_t - \eta \nabla J(w_t)$$

where  $\eta$  is the **learning rate**.

- **Types of Gradient Descent:**
    - **Batch:** Full dataset for stable but slow updates.
    - **Stochastic (SGD):** One data point for fast, noisy updates.
    - **Mini-Batch:** Small batches, balancing speed and stability.

## Problems with Gradient Descent

- **High Variability (SGD):** Quick in steep directions but slow in shallow ones, causing jitter and slow progress.
  - **Local Minima and Saddle Points:** Risk of **sub-optimal solutions** or long convergence times in flat regions.
  - **Noisy Updates:** Using individual points or mini-batches introduces noise, affecting stable convergence.



**Figure 4: SGD Variability (CS231n, Stanford)**

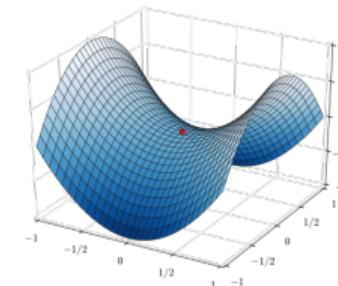


Figure 5: Saddle Point. Source: Wikipedia

## 1 Optimization

## 2 The Loss Surface

## 3 Gradient Descent

## 4 Momentum

First Moment (Momentum)

Second Moment (Variance)

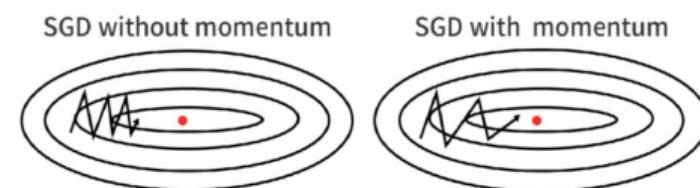
Adam: Adaptive Moment Estimation

## 5 Newton's optimization Method

## 6 References

# Problem Definition

- **Objective:** Enhance the vanilla Gradient Descent algorithm to improve convergence and stability.
- **Challenges:**
  - Selecting an appropriate learning rate is crucial to avoid slow convergence and getting stuck in local minima.
- **Proposed Solution:**
  - Instead of testing multiple learning rates, incorporate **Momentum** to adaptively adjust the learning rate based on oscillations:
    - Increase steps in stable directions.
    - Decrease steps in oscillating directions.



**Figure 6:** Momentum smooths oscillations and accelerates progress. Source: Papers with Code

# Introduction to Momentum in Optimization

- **Origin of Momentum:**

- Inspired by Newtonian physics, momentum in optimization uses **the concept of velocity in motion**, accumulating gradient history to smooth the learning trajectory, akin to an object moving based on past inertia.
- Initially introduced to tackle challenges in gradient descent, where **inconsistent gradients or noisy updates** lead to erratic and slow convergence.

- **Purpose of Momentum:**

- **Dampens Oscillations:** Utilizes prior gradients to minimize oscillations along steep or erratic regions, resulting in a smoother and more stable path.
- **Speeds Up Convergence:** Particularly effective in narrow valleys or flat regions, where standard gradient descent may struggle or oscillate, causing slow progress.

## 1 Optimization

## 2 The Loss Surface

## 3 Gradient Descent

## 4 Momentum

First Moment (Momentum)

Second Moment (Variance)

Adam: Adaptive Moment Estimation

## 5 Newton's optimization Method

## 6 References

## First Moment (Momentum)

- **Definition:** The first moment,  $m_t$ , represents a moving average of past gradients. It builds "velocity" that propels learning in a consistent direction.
  - **Update Rule:**

$$m_{t+1} = \beta_1 m_t + (1 - \beta_1) \nabla_w J(w_t)$$

$$w_{t+1} = w_t - \eta m_{t+1}$$

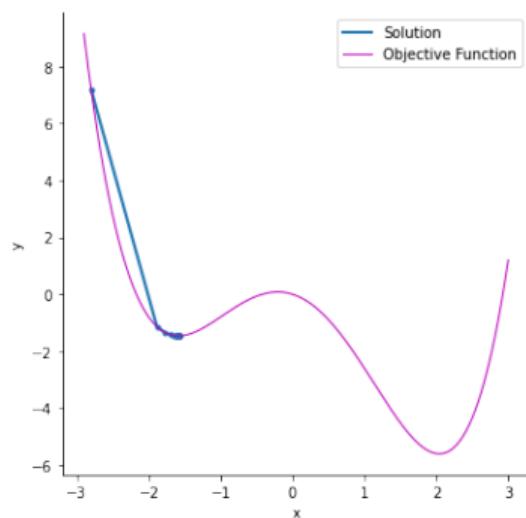
where:

- $\beta_1$ : Decay rate, usually 0.9 or 0.99, which controls the weight of past gradients.
  - $\eta$ : Learning rate.

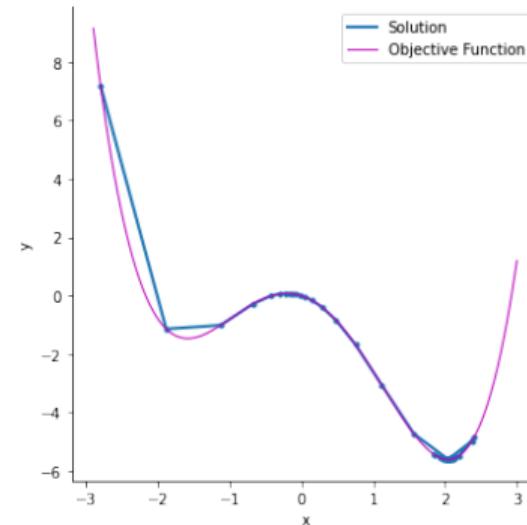
### • Why Use First Momentum?

- Inspired by the idea of rolling momentum, it smooths and accelerates learning by sustaining direction from prior gradients.
  - This type of momentum is ideal for traversing narrow valleys or regions where standard gradient descent would oscillate.

## Example of First Moment



**Figure 7:** Stochastic gradient descent without momentum stops at a local minimum. Source: Akash Ajagekar (SYSEN 6800 Fall 2021)



**Figure 8:** Stochastic gradient descent with momentum stops at the global minimum.  
Source: Akash Ajagekar (SYSEN 6800 Fall 2021)

## 1 Optimization

## 2 The Loss Surface

## 3 Gradient Descent

## 4 Momentum

First Moment (Momentum)

**Second Moment (Variance)**

Adam: Adaptive Moment Estimation

## 5 Newton's optimization Method

## 6 References

## Second Moment (Variance)

- **Definition:** The second moment,  $v_t$ , represents the moving average of squared gradients. It measures the gradient magnitude over time.
- **Update Rule:**

$$v_{t+1} = \beta_2 v_t + (1 - \beta_2)(\nabla_w J(w_t))^2$$

$$w_{t+1} = w_t - \frac{\eta}{\sqrt{v_{t+1} + \epsilon}} m_{t+1}$$

where:

- $\beta_2$ : Decay rate for variance (usually 0.99 or 0.999).
- $\epsilon$ : Small constant to prevent division by zero.

- **Why Use Second Momentum?**

- Adjusts step size based on gradient magnitude, preventing large steps when gradients are large and accelerating learning when they are small.

# Moment Bias Correction

- **Problem:** When we start training, both  $m_t$  and  $v_t$  are initialized to zero, causing their estimates to be **biased toward zero in the early steps**, especially when gradients are small.
- **Solution:** We use bias-corrected versions of  $m_t$  and  $v_t$  to address this:

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

- These corrections compensate for the bias by scaling  $m_t$  and  $v_t$  upward, especially in the early steps when  $t$  is small, ensuring more accurate estimates of the moments.

## 1 Optimization

## 2 The Loss Surface

## 3 Gradient Descent

## 4 Momentum

First Moment (Momentum)

Second Moment (Variance)

**Adam: Adaptive Moment Estimation**

## 5 Newton's optimization Method

## 6 References

# Introduction to Adam Optimizer

- **Origin and Purpose:**

- Proposed in 2014 by Diederik Kingma and Jimmy Ba, Adam (Adaptive Moment Estimation) addresses key limitations in earlier optimization methods by combining aspects of **momentum** and **adaptive learning rates**.
- Adam is designed to handle sparse gradients and noisy updates by adjusting the learning rate for each parameter based on historical gradients.

- **Core Idea:**

- Adam optimizes by maintaining two moving averages — the **first moment (mean of gradients)** and the **second moment (variance of gradients)** — allowing it to **adapt learning rates for each parameter individually**.

# Adam's Adaptive Learning Rate Mechanism

- **Why Adaptive Rates?**

- Unlike traditional SGD, Adam adapts the learning rate for **each parameter** based on recent gradient magnitudes.
- **Large gradients** lead to **reduced** update sizes, while **smaller gradients** allow **larger** updates, balancing convergence speed and stability.

- **Moment Tracking**

- The **first moment** ( $m_t$ ) tracks the mean of gradients to provide momentum.
- The **second moment** ( $v_t$ ) tracks squared gradients, enabling Adam to normalize updates and prevent sudden changes in direction.

# Mathematical Formulation of Adam

- **Adam Update Rules:**

- First moment estimate:

$$m_{t+1} = \beta_1 m_t + (1 - \beta_1) \nabla_w J(w_t)$$

- Second moment estimate:

$$v_{t+1} = \beta_2 v_t + (1 - \beta_2) (\nabla_w J(w_t))^2$$

- Bias-corrected moments to address initialization bias:

$$\hat{m}_{t+1} = \frac{m_{t+1}}{1 - \beta_1^{t+1}}, \quad \hat{v}_{t+1} = \frac{v_{t+1}}{1 - \beta_2^{t+1}}$$

- Update step for parameter  $w_t$ :

$$w_{t+1} = w_t - \eta \frac{\hat{m}_{t+1}}{\sqrt{\hat{v}_{t+1}} + \epsilon}$$

# Adam Pseudo-code

---

**Algorithm 1:** Adam, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation.  $g_t^2$  indicates the elementwise square  $g_t \odot g_t$ . Good default settings for the tested machine learning problems are  $\alpha = 0.001$ ,  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$  and  $\epsilon = 10^{-8}$ . All operations on vectors are element-wise. With  $\beta_1^t$  and  $\beta_2^t$  we denote  $\beta_1$  and  $\beta_2$  to the power  $t$ .

---

**Require:**  $\alpha$ : Stepsize  
**Require:**  $\beta_1, \beta_2 \in [0, 1]$ : Exponential decay rates for the moment estimates  
**Require:**  $f(\theta)$ : Stochastic objective function with parameters  $\theta$   
**Require:**  $\theta_0$ : Initial parameter vector

$m_0 \leftarrow 0$  (Initialize 1<sup>st</sup> moment vector)  
 $v_0 \leftarrow 0$  (Initialize 2<sup>nd</sup> moment vector)  
 $t \leftarrow 0$  (Initialize timestep)  
**while**  $\theta_t$  not converged **do**

- $t \leftarrow t + 1$
- $g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$  (Get gradients w.r.t. stochastic objective at timestep  $t$ )
- $m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$  (Update biased first moment estimate)
- $v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$  (Update biased second raw moment estimate)
- $\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$  (Compute bias-corrected first moment estimate)
- $\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$  (Compute bias-corrected second raw moment estimate)
- $\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$  (Update parameters)

**end while**  
**return**  $\theta_t$  (Resulting parameters)

---

Figure 9: Adam Pseudo-code. Source: kingma2014adam

# Adam Visualization

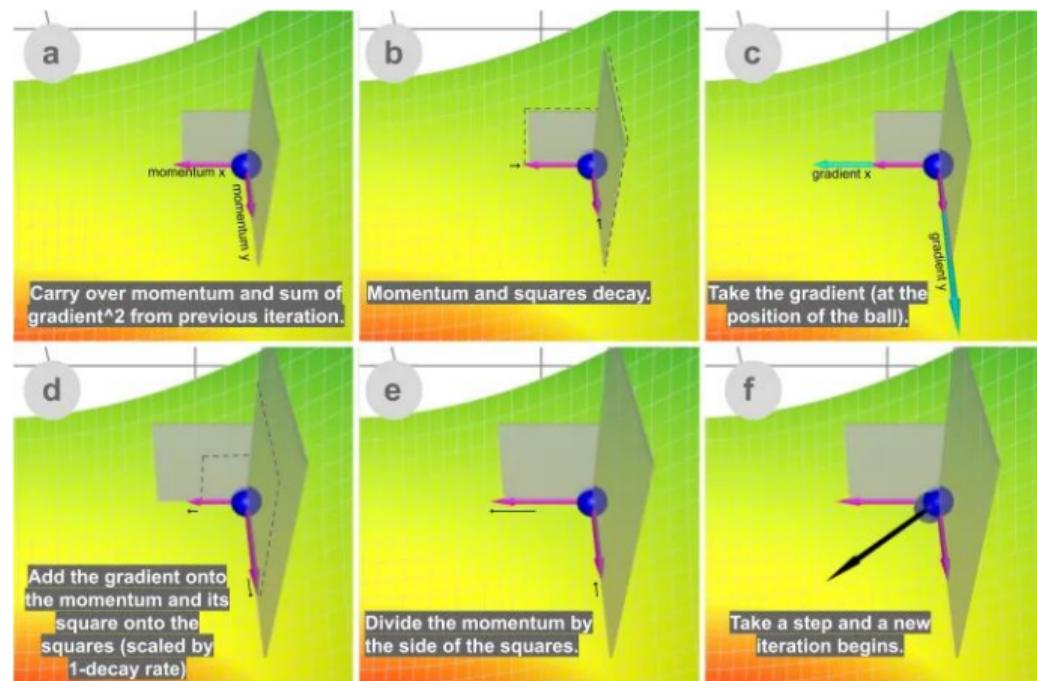
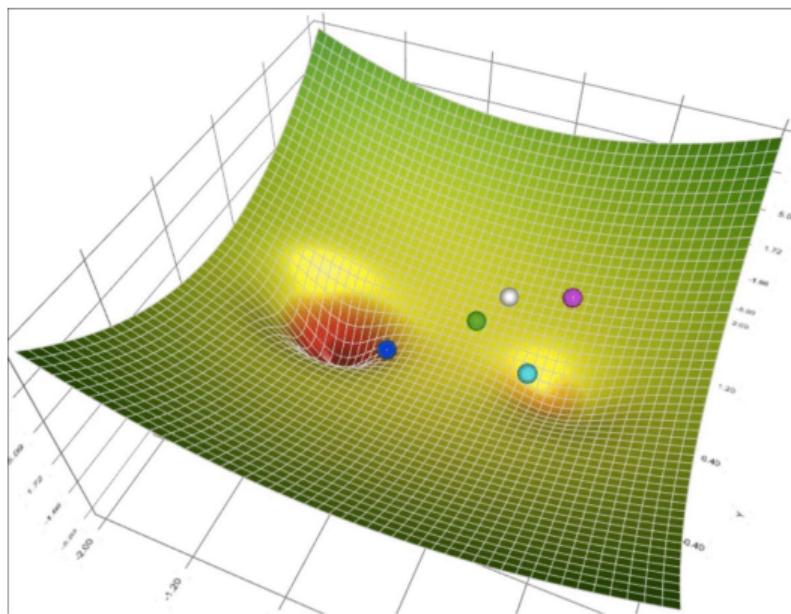


Figure 10: Step-by-step illustration of Adam descent. Source: Towards Data Science

# Comparison of Momentum Methods



**Figure 11:** Comparison of 5 gradient descent methods on a surface: gradient descent (cyan), momentum (magenta), AdaGrad (white), RMSProp (green), Adam (blue). Left well is the global minimum; right well is a local minimum. Source: Towards Data Science

## 1 Optimization

## 2 The Loss Surface

## 3 Gradient Descent

## 4 Momentum

## 5 Newton's optimization Method

Newton's Method

## 6 References

## 1 Optimization

## 2 The Loss Surface

## 3 Gradient Descent

## 4 Momentum

## 5 Newton's optimization Method

### Newton's Method

## 6 References

## Newton's Method

- Newton method is originally intended to **find the root(s)** of an equation.
  - **Example:** for the equation  $x^2 - 1 = 0$ , we can find the roots by decomposing  $(x - 1)(x + 1) = 0$  which gives  $x = 1, x = -1$
  - **But, what about complex equations?**
    - We can use **numerical method** to find the root of an equation, one of them is by using **Newton's method**

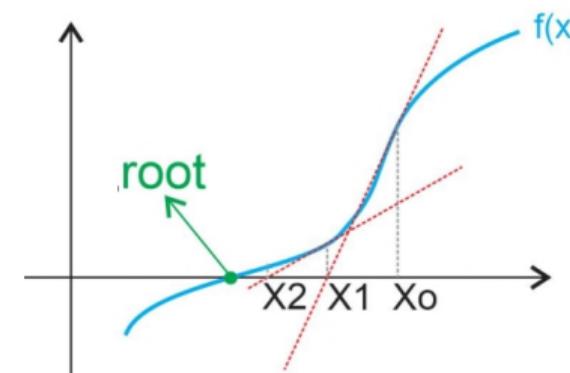
### Definition

- **Objective:** Derive Newton's method by finding the tangent line of  $f(x)$  at  $x_0$ .
  - **Tangent Line Equation:** Given a point  $x_0$  where  $f(x_0) \neq 0$ , the tangent line at  $x_0$  is:

$$y = mx_0 + c$$

- **Gradient:** The slope  $m$  matches the derivative of  $f(x)$  at  $x_0$ :

$$m = f'(x_0)$$



**Figure 12:** Finding root location by using Newton's method. Source: Ardian Umam's Blog

## Formulating the Tangent Line

- **Finding  $c$ :** Substitute  $(x_0, f(x_0))$  into  $y = mx + c$ , where  $y = f(x_0)$  and  $m = f'(x_0)$ :

$$f(x_0) = f'(x_0)x_0 + c \Rightarrow c = f(x_0) - f'(x_0)x_0$$

- **Tangent Line Equation:** Substitute  $m = f'(x_0)$  and  $c$  back:

$$v \equiv f'(x_0)x + f(x_0) - f'(x_0)x_0$$

- Simplify to get:

$$y = f(x_0) + f'(x_0)(x - x_0)$$

## Newton's Iterative Step

- To approximate the root, set  $y=0$  in the tangent equation:

$$0 = f(x_0) + f'(x_0)(x_1 - x_0)$$

- Rearrange to solve for  $x_1$ :

$$x_1 = x_0 - \frac{f(x_0)}{f'(x_0)}$$

- **Iteration:** Repeat this step to approximate the root.

# Newton's Method for Optimization

- Newton's method for finding roots is based on a first-order approximation (tangent line).
  - For optimization, we use a second-order Taylor approximation to find the minimum.
  - **Second-order Taylor expansion** of  $f(x)$  around  $x = x_0$ :

$$f(x) \approx f(x_0) + f'(x_0)(x - x_0) + \frac{f''(x_0)}{2}(x - x_0)^2$$

- Rearranged for minimal value location:

$$f(x) \approx \frac{1}{2}f''(x_0)x^2 + [f'(x_0) - f''(x_0)x_0]x + [f(x_0) - f'(x_0)x_0 + \frac{1}{2}f''(x_0)x_0^2]$$

# Deriving the Update Formula for Minimization

- To locate the minimum, take the derivative with respect to  $x$  and set it to zero:

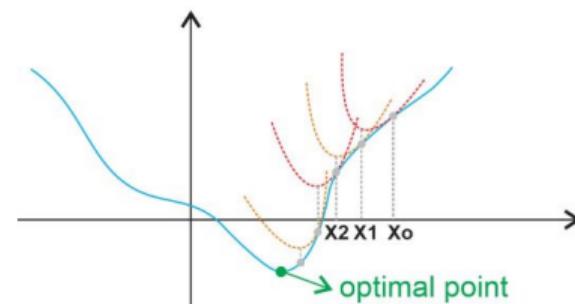
$$\frac{d}{dx}f(x) \approx f''(x_0)x + [f'(x_0) - f''(x_0)x_0] = 0$$

- Solving for  $x$  yields:

$$x = x_0 - \frac{f'(x_0)}{f''(x_0)}$$

- This is the update step for Newton's method in optimization, guiding us to the minimum. The general update rule is:

$$x_{t+1} = x_t - H^{-1} \nabla_x f(x_t)$$



**Figure 13:** Finding root using Taylor's expansion and Newton's method. Source: Ardian Umam's Blog

# Hessian Matrix and Newton's Method for Optimization

- The Hessian matrix,  $H(\theta)$ , is a square matrix of second-order partial derivatives of a scalar-valued function  $f(\theta)$ :

$$H(\theta) = \begin{bmatrix} \frac{\partial^2 f}{\partial \theta_1^2} & \frac{\partial^2 f}{\partial \theta_1 \partial \theta_2} & \cdots & \frac{\partial^2 f}{\partial \theta_1 \partial \theta_n} \\ \frac{\partial^2 f}{\partial \theta_2 \partial \theta_1} & \frac{\partial^2 f}{\partial \theta_2^2} & \cdots & \frac{\partial^2 f}{\partial \theta_2 \partial \theta_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial \theta_n \partial \theta_1} & \frac{\partial^2 f}{\partial \theta_n \partial \theta_2} & \cdots & \frac{\partial^2 f}{\partial \theta_n^2} \end{bmatrix}$$

- In Newton's method for optimization the update rule for parameters  $\theta$  is:

$$\theta_{t+1} = \theta_t - H^{-1}(\theta_t) \nabla f(\theta_t)$$

- Example:

$$f(\theta_1, \theta_2) = \theta_1^2 + 2\theta_2^2$$

$$\nabla f = \begin{bmatrix} \frac{\partial f}{\partial \theta_1} \\ \frac{\partial f}{\partial \theta_2} \end{bmatrix} = \begin{bmatrix} 2\theta_1 \\ 4\theta_2 \end{bmatrix}$$

$$H(\theta) = \begin{bmatrix} \frac{\partial^2 f}{\partial \theta_1^2} & \frac{\partial^2 f}{\partial \theta_1 \partial \theta_2} \\ \frac{\partial^2 f}{\partial \theta_2 \partial \theta_1} & \frac{\partial^2 f}{\partial \theta_2^2} \end{bmatrix} = \begin{bmatrix} 2 & 0 \\ 0 & 4 \end{bmatrix}$$

$$H^{-1} = \begin{bmatrix} 1/2 & 0 \\ 0 & 1/4 \end{bmatrix}$$

$$\theta_{t+1} = \theta_t - H^{-1} \nabla f(\theta_t)$$

## Newton's Method: Advantages and Disadvantages

- Newton's method offers various benefits but also has limitations, especially in large-scale machine learning. Below is a summary:

Advantages	Disadvantages
<b>Faster Convergence</b> Quadratic convergence enables reaching minima faster in convex problems.	<b>Computationally Expensive</b> Requires Hessian calculation, making it costly in high-dimensional models.
<b>Adaptive Step Sizes</b> Curvature-based step adjustment avoids slow progress in shallow regions.	<b>Memory Intensive</b> Storing the Hessian matrix is memory-intensive for models with millions of parameters.
<b>Reduced Oscillations</b> Curvature information stabilizes paths in oscillatory regions.	<b>Convergence Challenges</b> May converge to saddle points in non-convex functions common in machine learning.

Table 1: Advantages and Disadvantages of Newton's Method

## ① Optimization

## ② The Loss Surface

## ③ Gradient Descent

## ④ Momentum

## ⑤ Newton's optimization Method

## ⑥ References

## Contribution

- These slides were prepared with contributions from:
    - Alireza Sabounchi
    - Sina Daneshgar

