

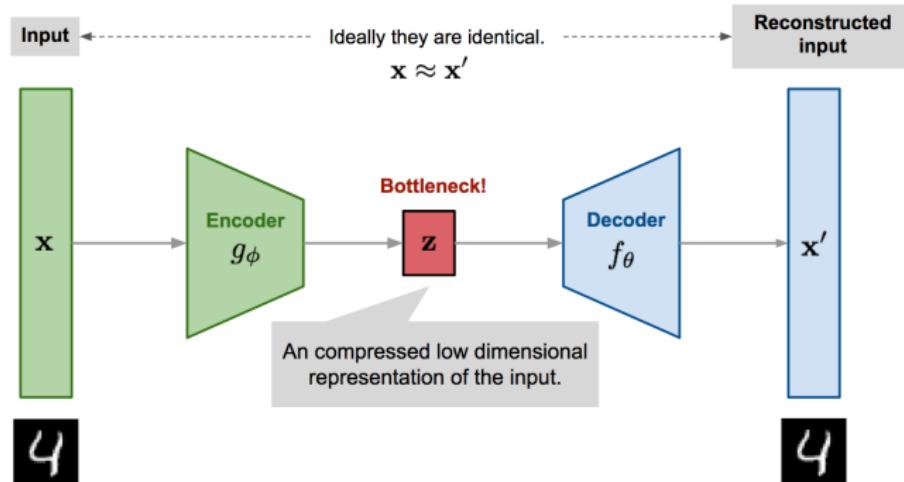
# AutoEncoders

ML Instruction Team, Fall 2022

CE Department  
Sharif University of Technology

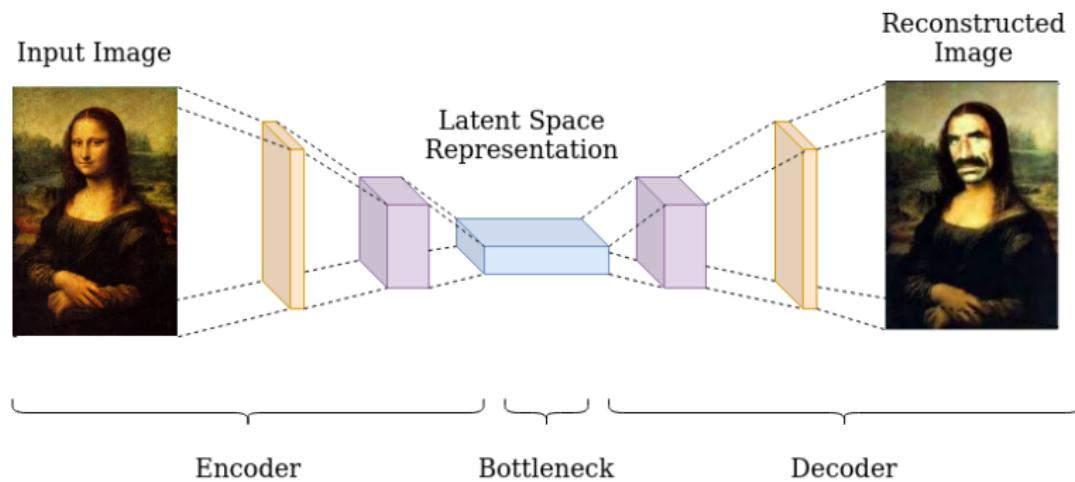
# Autoencoders

- Autoencoders are artificial neural networks capable of learning dense representations of the input data, called *latent representations* or *codings*, without any supervision (i.e., the training set is unlabeled).
- Their job is to take an input  $X$  and predict  $X$ . To make this non-trivial, we need to add a **bottleneck layer** whose dimension is much smaller than the input.



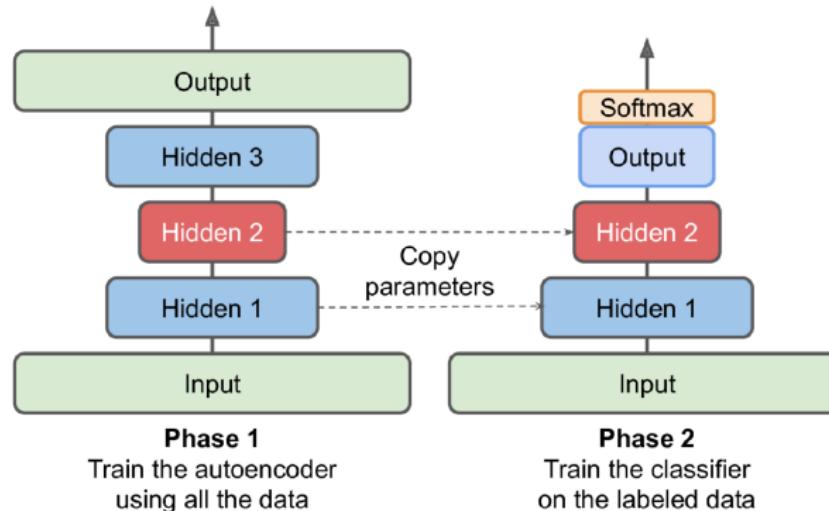
# Autoencoders: Structure

- Encoder: compress input into a latent-space of usually smaller dimension.  $h = f(x)$
- Decoder: reconstruct input from the latent space.  $r = g(f(x))$  with  $r$  as close to  $x$  as possible



# Autoencoders: Applications

- Autoencoders can
  - ▶ act as feature detectors
  - ▶ be used for unsupervised pretraining of deep neural networks



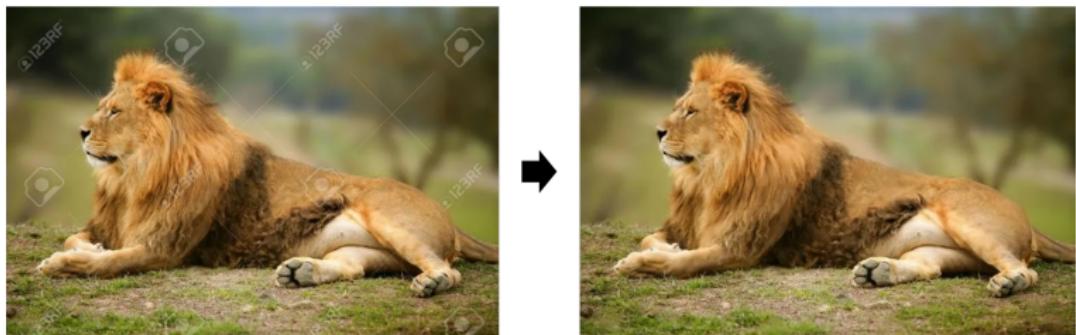
# Autoencoders: Applications

- Autoencoders can be used as generative models. (will be discussed more later in this chapter.)



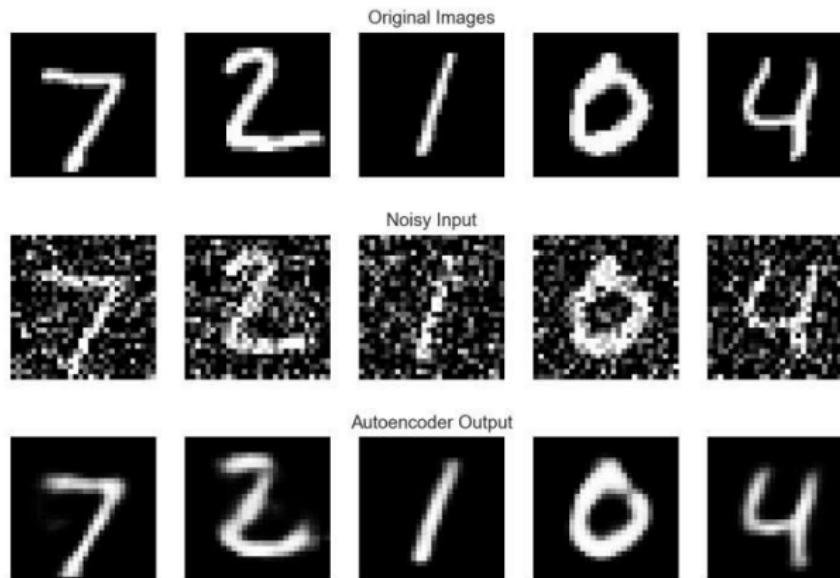
# Autoencoders: Applications

- Watermark removal



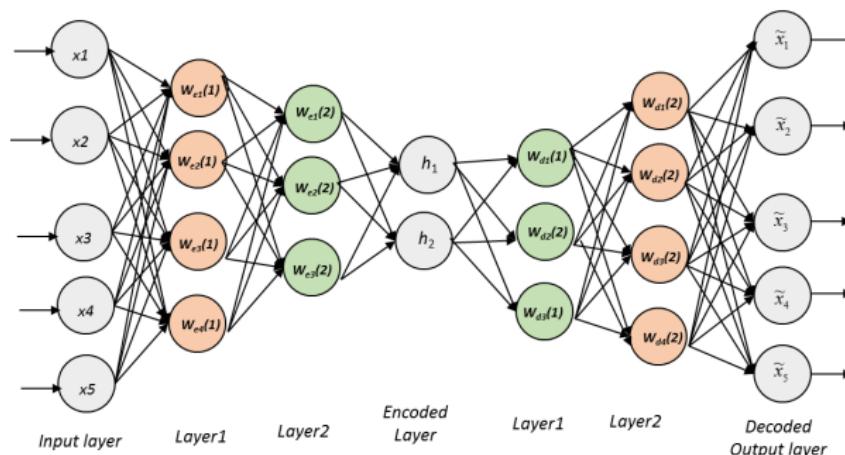
# Autoencoders: Applications

## ■ Noise reduction



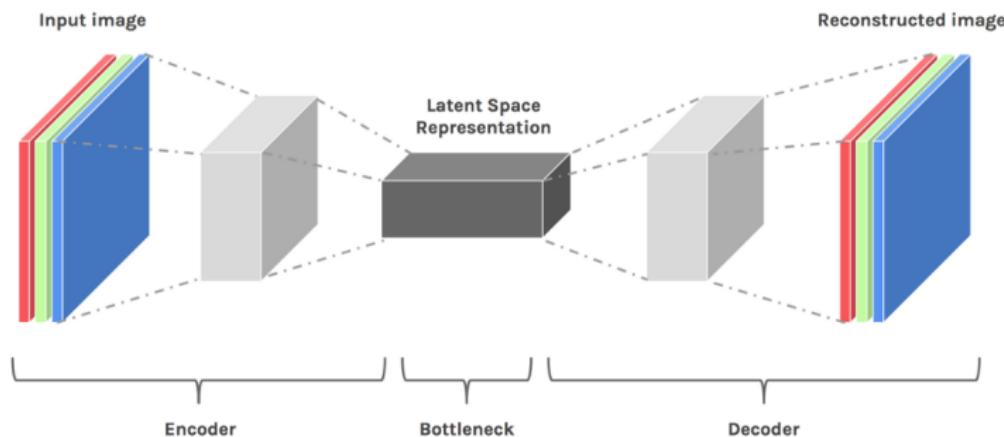
# Stacked Autoencoders

- Autoencoders can have multiple hidden layers. In this case they are called *stacked autoencoders* (or *deep autoencoders*).
- Adding more layers helps the autoencoder learn more complex codings, but be careful about **overfitting!**



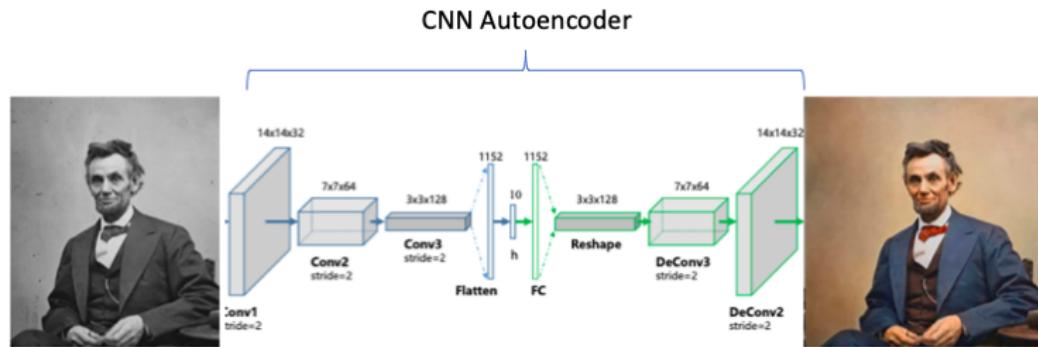
# Autoencoders & Images

**Are normal Autoencoders suitable for working with images?**



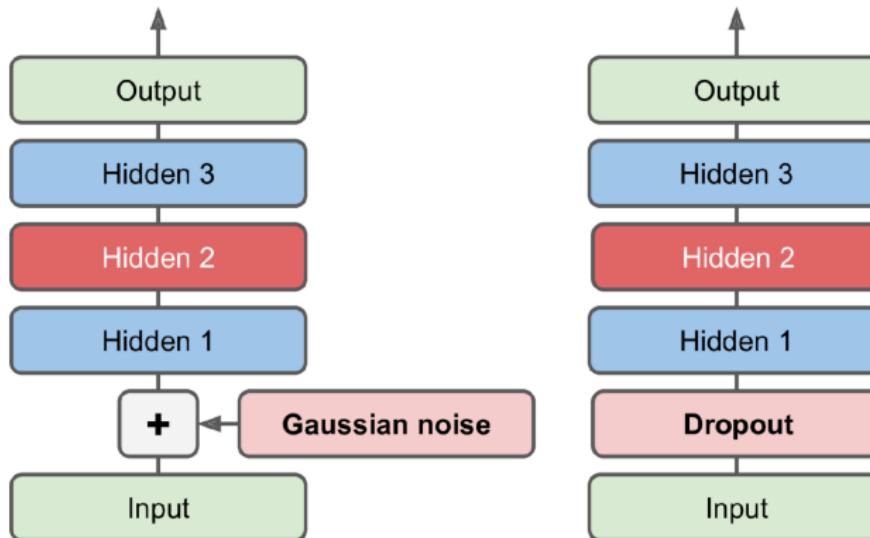
# Autoencoders & Images

## ■ Convolutional Autoencoders



# Denoising Autoencoders

- Another way to force the autoencoder to learn useful features is to add noise to its inputs.
- Denoising autoencoders train to minimize the loss between  $x$  and  $g(f(x + w))$ , where  $w$  is random noise.
- Denoising autoencoders, with Gaussian noise (left) or dropout (right):



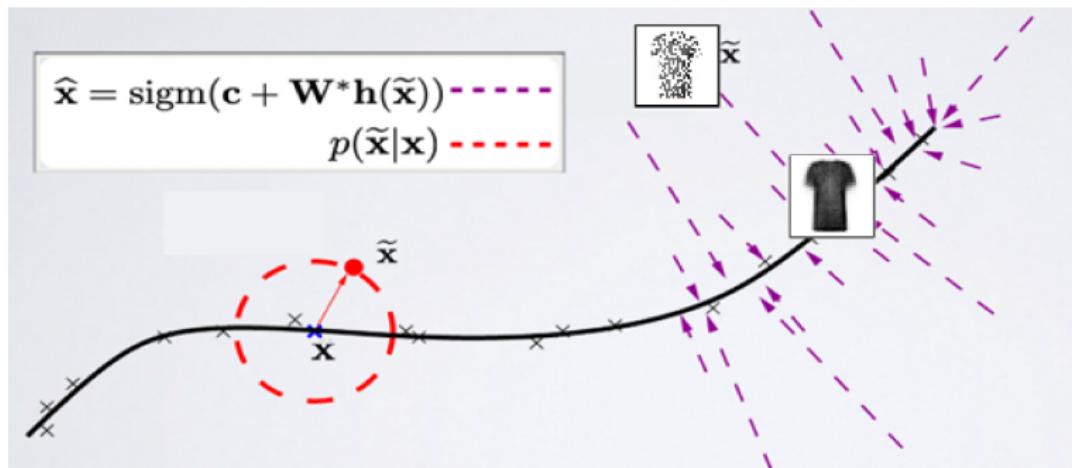
# Denoising Autoencoders

- A few noisy images (with half the pixels turned off), and the images reconstructed by the dropout-based denoising autoencoder. Notice how the autoencoder guesses details that are actually not in the input, such as the top of the white shirt (bottom row, fourth image).



# Denoising Autoencoders

- Intuitively, a denoising autoencoder learns a projection from a neighborhood of our training data back onto the training data.



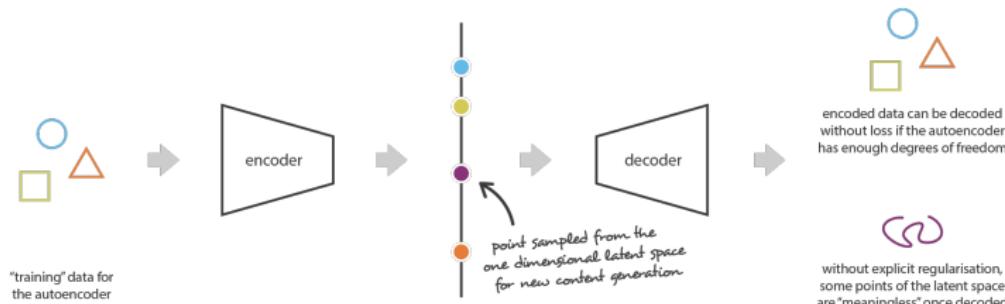
# Autoencoder Generative Models

**How can we generate NEW data with Autoencoders??**

hint: Autoencoder learns the feature space!

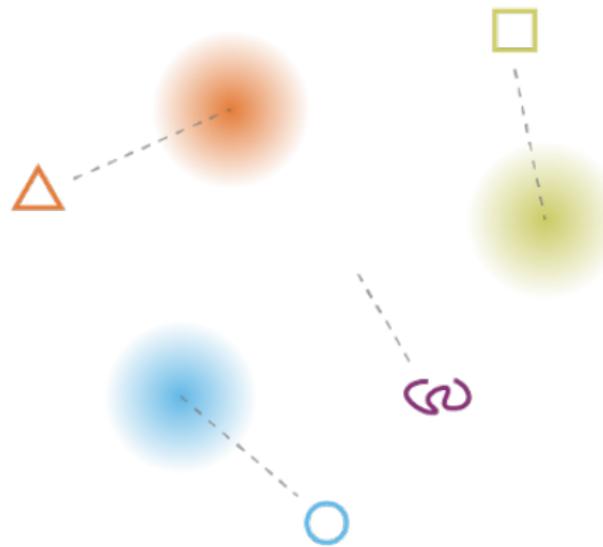
# Walking through an example

- We want to reconstruct some shapes.



# Walking through an example

- Not all of the points in latent space have meaningful reconstructions.



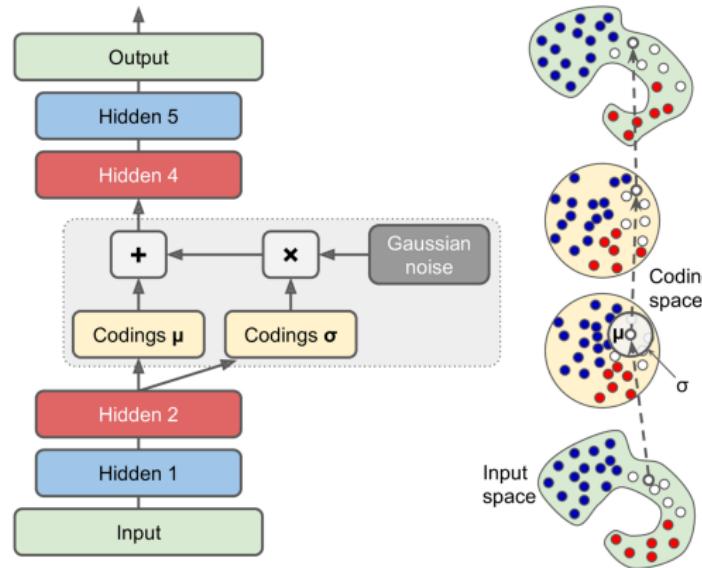
# Walking through an example

- What we want is something like the following picture. So that with sampling from the latent space, we can generate new shapes.

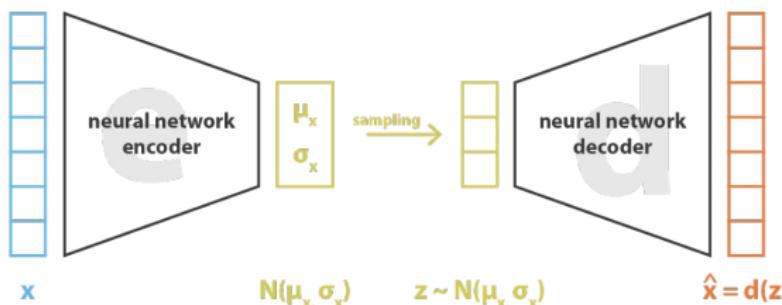


# Variational Autoencoders

- instead of directly producing a coding for a given input, the encoder produces a mean coding  $\mu$  and a standard deviation  $\sigma$ . The actual coding is then sampled randomly from a Gaussian distribution with mean  $\mu$  and standard deviation  $\sigma$



# Variational Autoencoders



---


$$\text{loss} = \|x - \hat{x}\|^2 + \text{KL}[N(\mu_x, \sigma_x), N(0, I)] = \|x - d(z)\|^2 + \text{KL}[N(\mu_x, \sigma_x), N(0, I)]$$

# Test

- One
  - ▶ One
  - ▶ Two
  - ▶ Three
- For two-dimensional tensors, we have a corresponding sum with indices  $(a, b)$  for  $f$  and  $(i - a, j - b)$  for  $g$ , respectively:

$$(f * g)(i, j) = \sum_a \sum_b f(a, b)g(i - a, j - b)$$

- It is given by,

$$w_{t+1} = w_t - \left( \alpha_t / \sqrt(v_t) + e \right) * (\delta L / \delta w_t)$$

where,

$$v_t = \beta * v_t + (1 - \beta) * (\delta L / \delta w_t)^2$$

# Image References

- <https://lilianweng.github.io/posts/2018-08-12-vae/>
- <https://emkadem.y.medium.com/1-first-step-to-generative-deep-learning-with-autoencoders-22bd41e56d18>
- Aurelien Geron. 2019. Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems (2nd. ed.). O'Reilly Media, Inc.
- <https://github.com/wojciechmo/vae>
- <https://ai.googleblog.com/2017/08/making-visible-watermarks-more-effective.html>
- <https://medium.com/@harishr2301/denoising-autoencoders-996e866e5cd0>
- <https://subscription.packtpub.com/book/big-data-and-business-intelligence/9781787121089/4/ch04lv11sec51/setting-up-stacked-autoencoders>
- <https://www.analyticsvidhya.com/blog/2021/01/auto-encoders-for-computer-vision-an-endless-world-of-possibilities/>
- <https://towardsdatascience.com/convolutional-autoencoders-for-image-noise-reduction-32fce9fc1763>
- [https://ift6266h17.files.wordpress.com/2017/03/14\\_autoencoders.pdf](https://ift6266h17.files.wordpress.com/2017/03/14_autoencoders.pdf)
- <https://towardsdatascience.com/understanding-variational-autoencoders-vaes-f70510919f73>

# References

- Aurelien Geron. 2019. Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems (2nd. ed.). O'Reilly Media, Inc.
- [https://www.cs.toronto.edu/~rgrosse/courses/csc321\\_2017/slides/lec20.pdf](https://www.cs.toronto.edu/~rgrosse/courses/csc321_2017/slides/lec20.pdf)
- <https://www.math.purdue.edu/~buzzard/MA598-Spring2019/Lectures/Lec16%20-%20Autoencoders.pptx>
- [http://cs231n.stanford.edu/slides/2019/cs231n\\_2019\\_lecture11.pdf](http://cs231n.stanford.edu/slides/2019/cs231n_2019_lecture11.pdf)

**Thank You!**

**Any Question?**